

Understanding Energy Behaviors of Thread Management Constructs



Gustavo Pinto¹



Fernando Castor¹



David Liu²

{ghlp, castor}@cin.ufpe.br¹
davidL@binghamton.edu²



BINGHAMTON
UNIVERSITY
STATE UNIVERSITY OF NEW YORK

Motivation



- **First**, the proliferation of multicore CPUs
- **Second**, the prevalence of multi-threaded programs

The Problem

- **First**, more cores → more power consumed
- **Second**, little is known about **energy behaviors** of multi-threaded program on the **application and programming language** level

This Talk

1. **programming abstractions** of thread management
on energy efficiency
2. **programmer choices** of thread management on
energy efficiency

for Java multi-threaded programs

Thread management constructs

- **Explicit threading (the Thread-style):** Using the *java.lang.Thread* class
- **Thread pooling (the Executor-style):** Using the *java.util.concurrent.Executor* framework
- **Working Stealing (the ForkJoin-style):** Using the *java.util.concurrent.ForkJoin* framework

Benchmarks

- **Embarrassingly parallel:** spectralnorm, sunflow, n-queens
- **Leaning parallel:** xalan, knucleotide, tomcat
- **Leaning serial:** mandelbrot, largestImage
- **Embarrassingly serial:** h2

Benchmarks

- **Embarrassingly parallel:** spectralnorm, sunflow,
n-queens
- **Leaning parallel:** xalan, knucleotide, tomcat
- **Leaning serial:** mandelbrot, largestImage
- **Embarrassingly serial:** h2

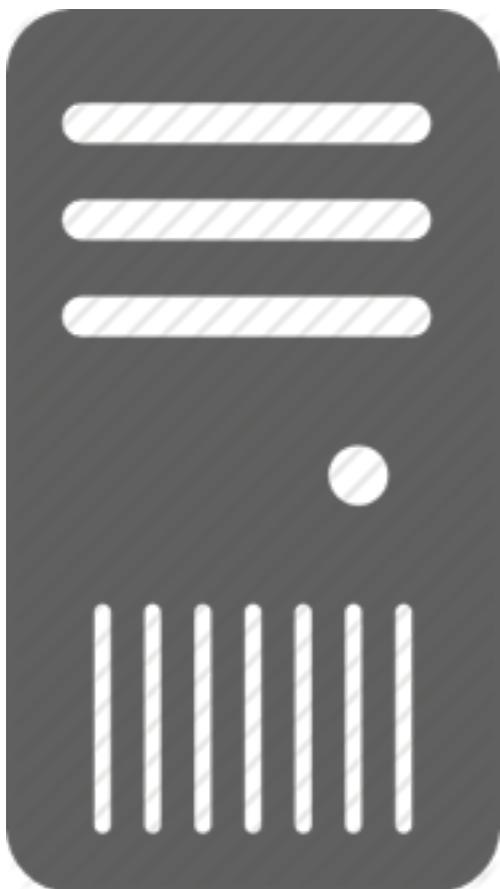


Micro-benchmarks



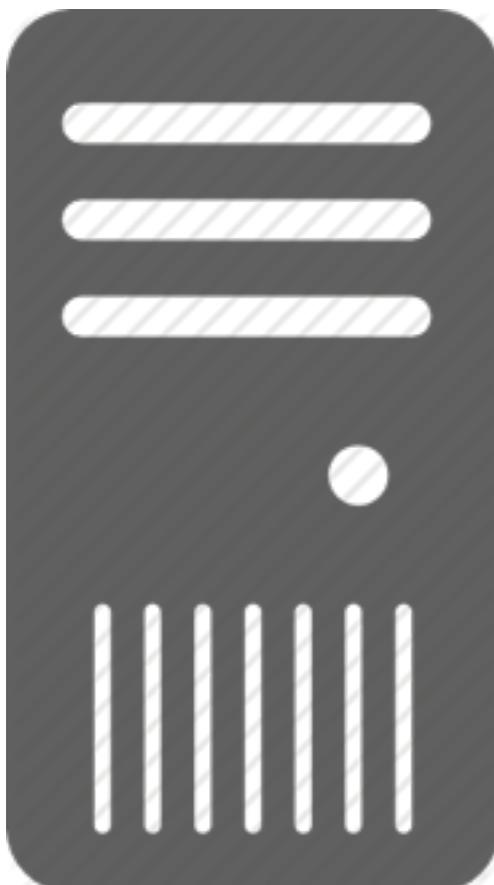
DaCapo benchmarks

Experimental Environment

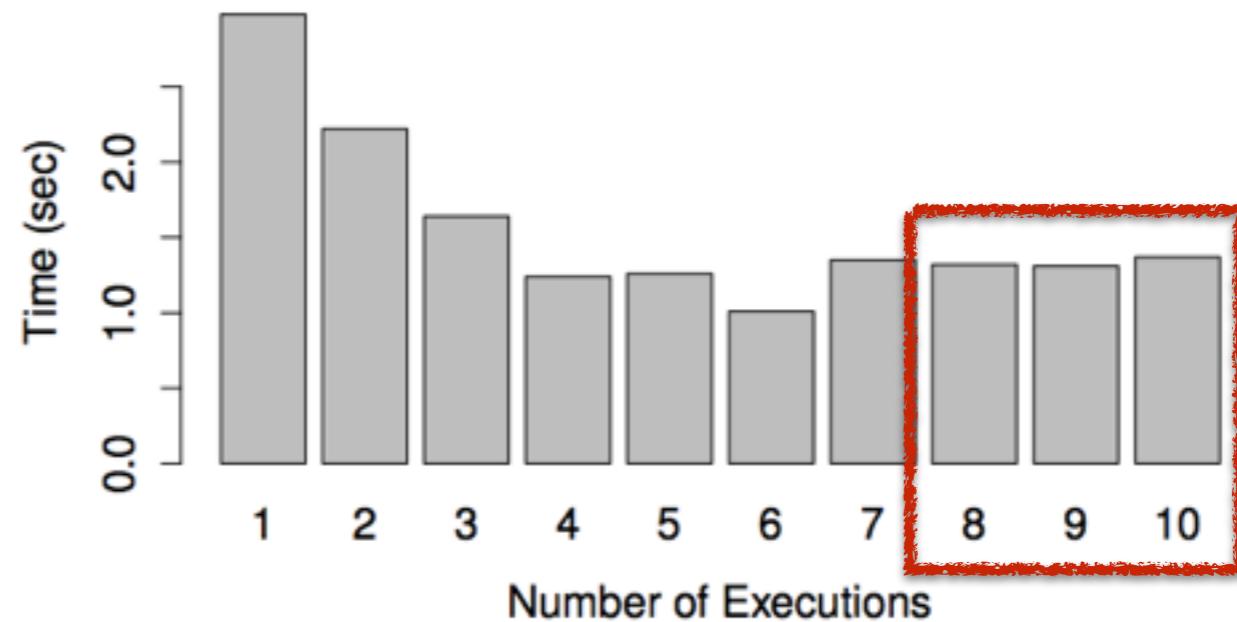


A 2×16-core AMD CPUs, running Debian Linux, 64GB of memory, JDK version 1.7.0 11, build 21, “ondemand” governor

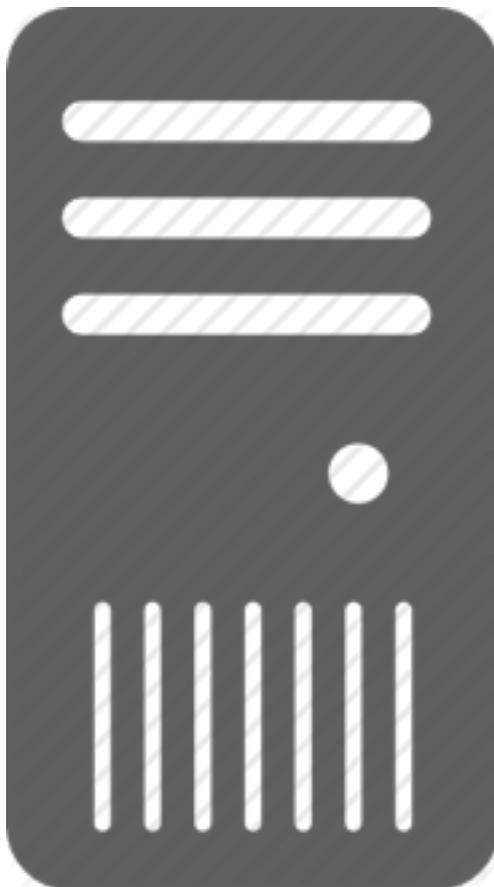
Experimental Environment



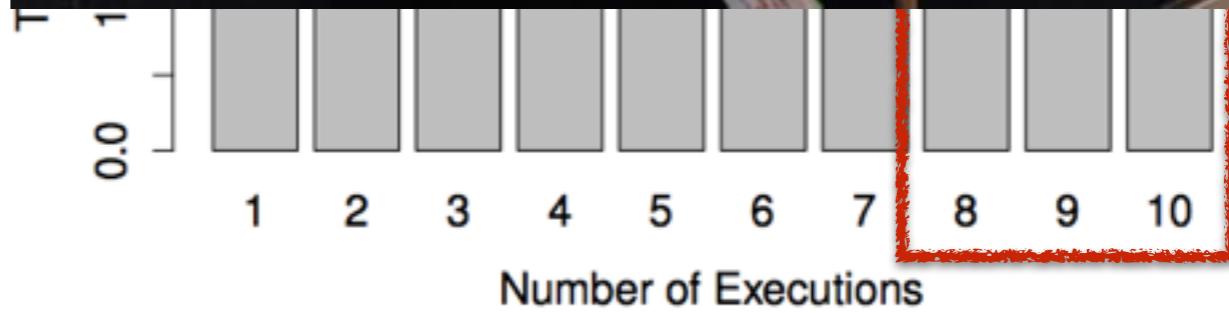
A 2×16-core AMD CPUs, running Debian Linux, 64GB of memory, JDK version 1.7.0 11, build 21, “ondemand” governor



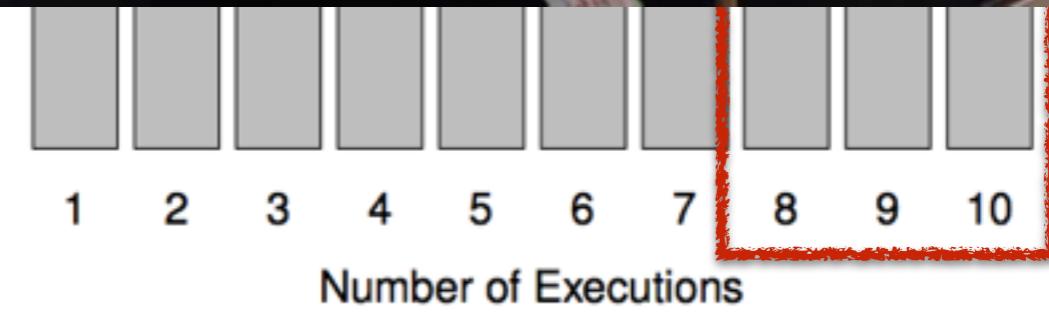
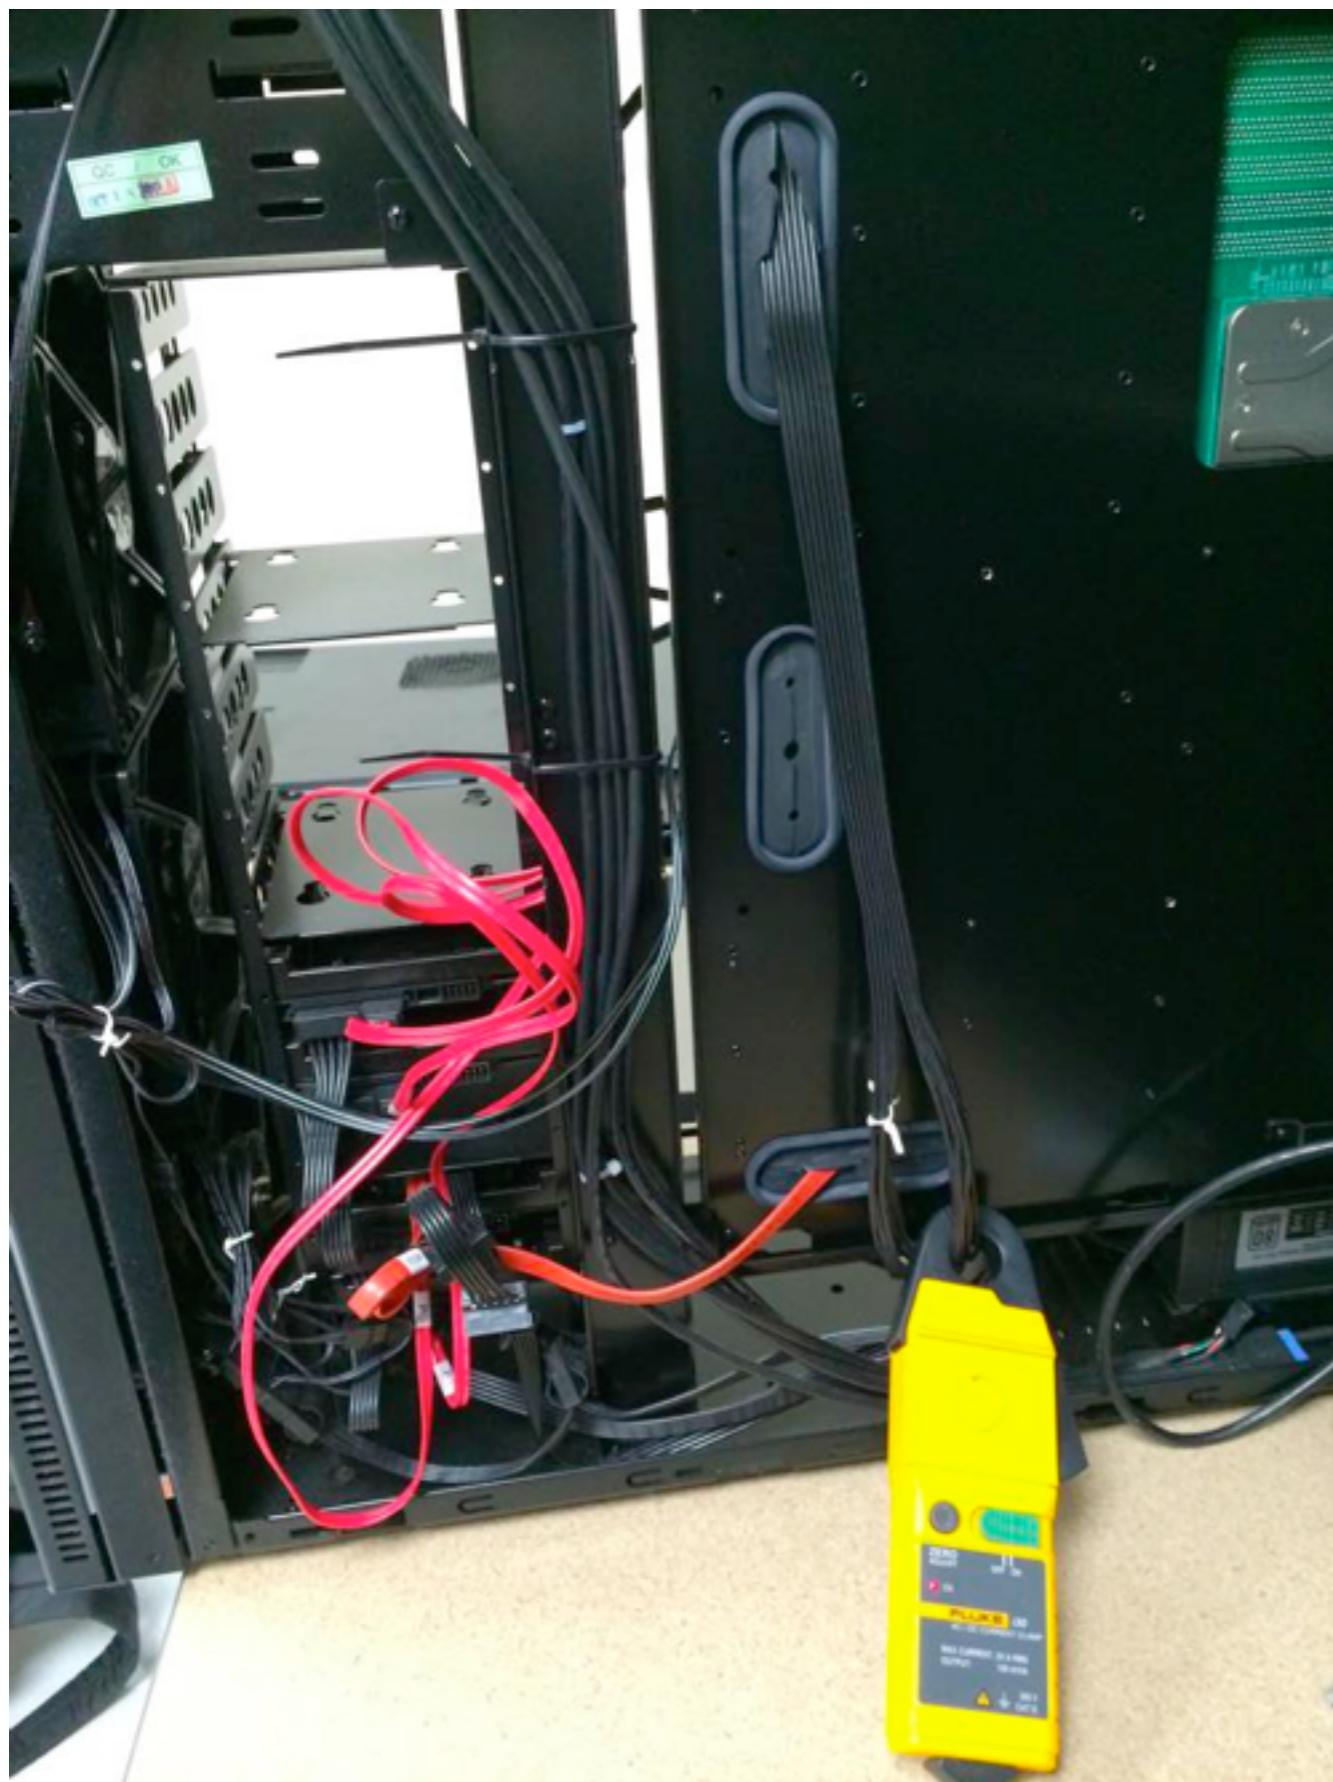
Experimental Environment



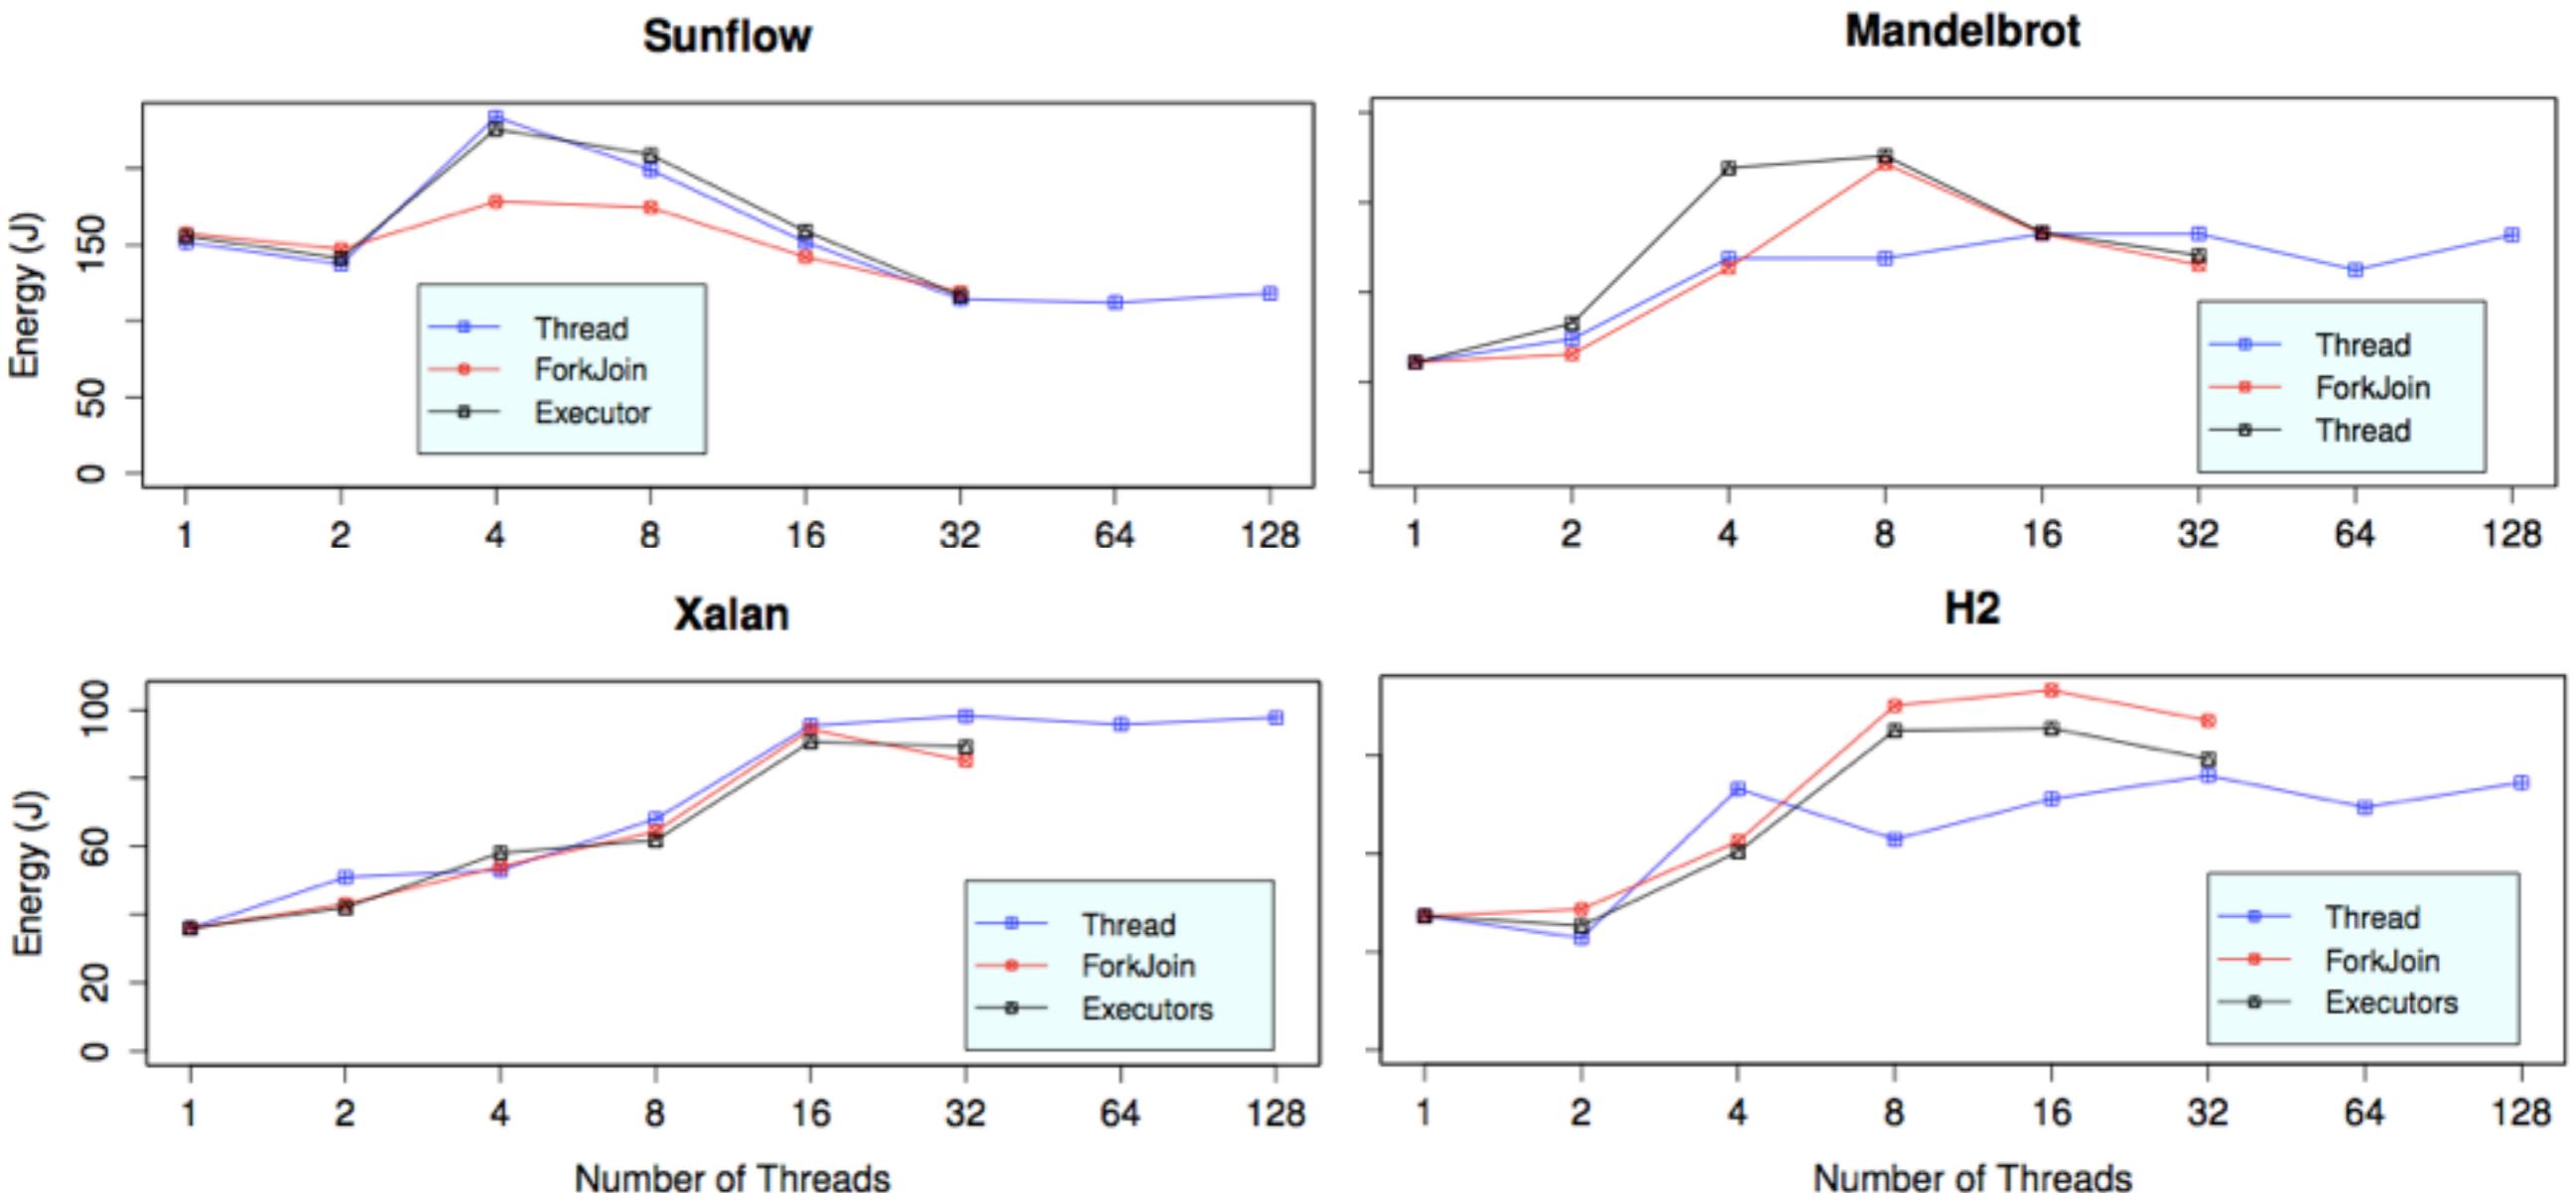
A 2×16-core
Linux, 64GB
11, build 21



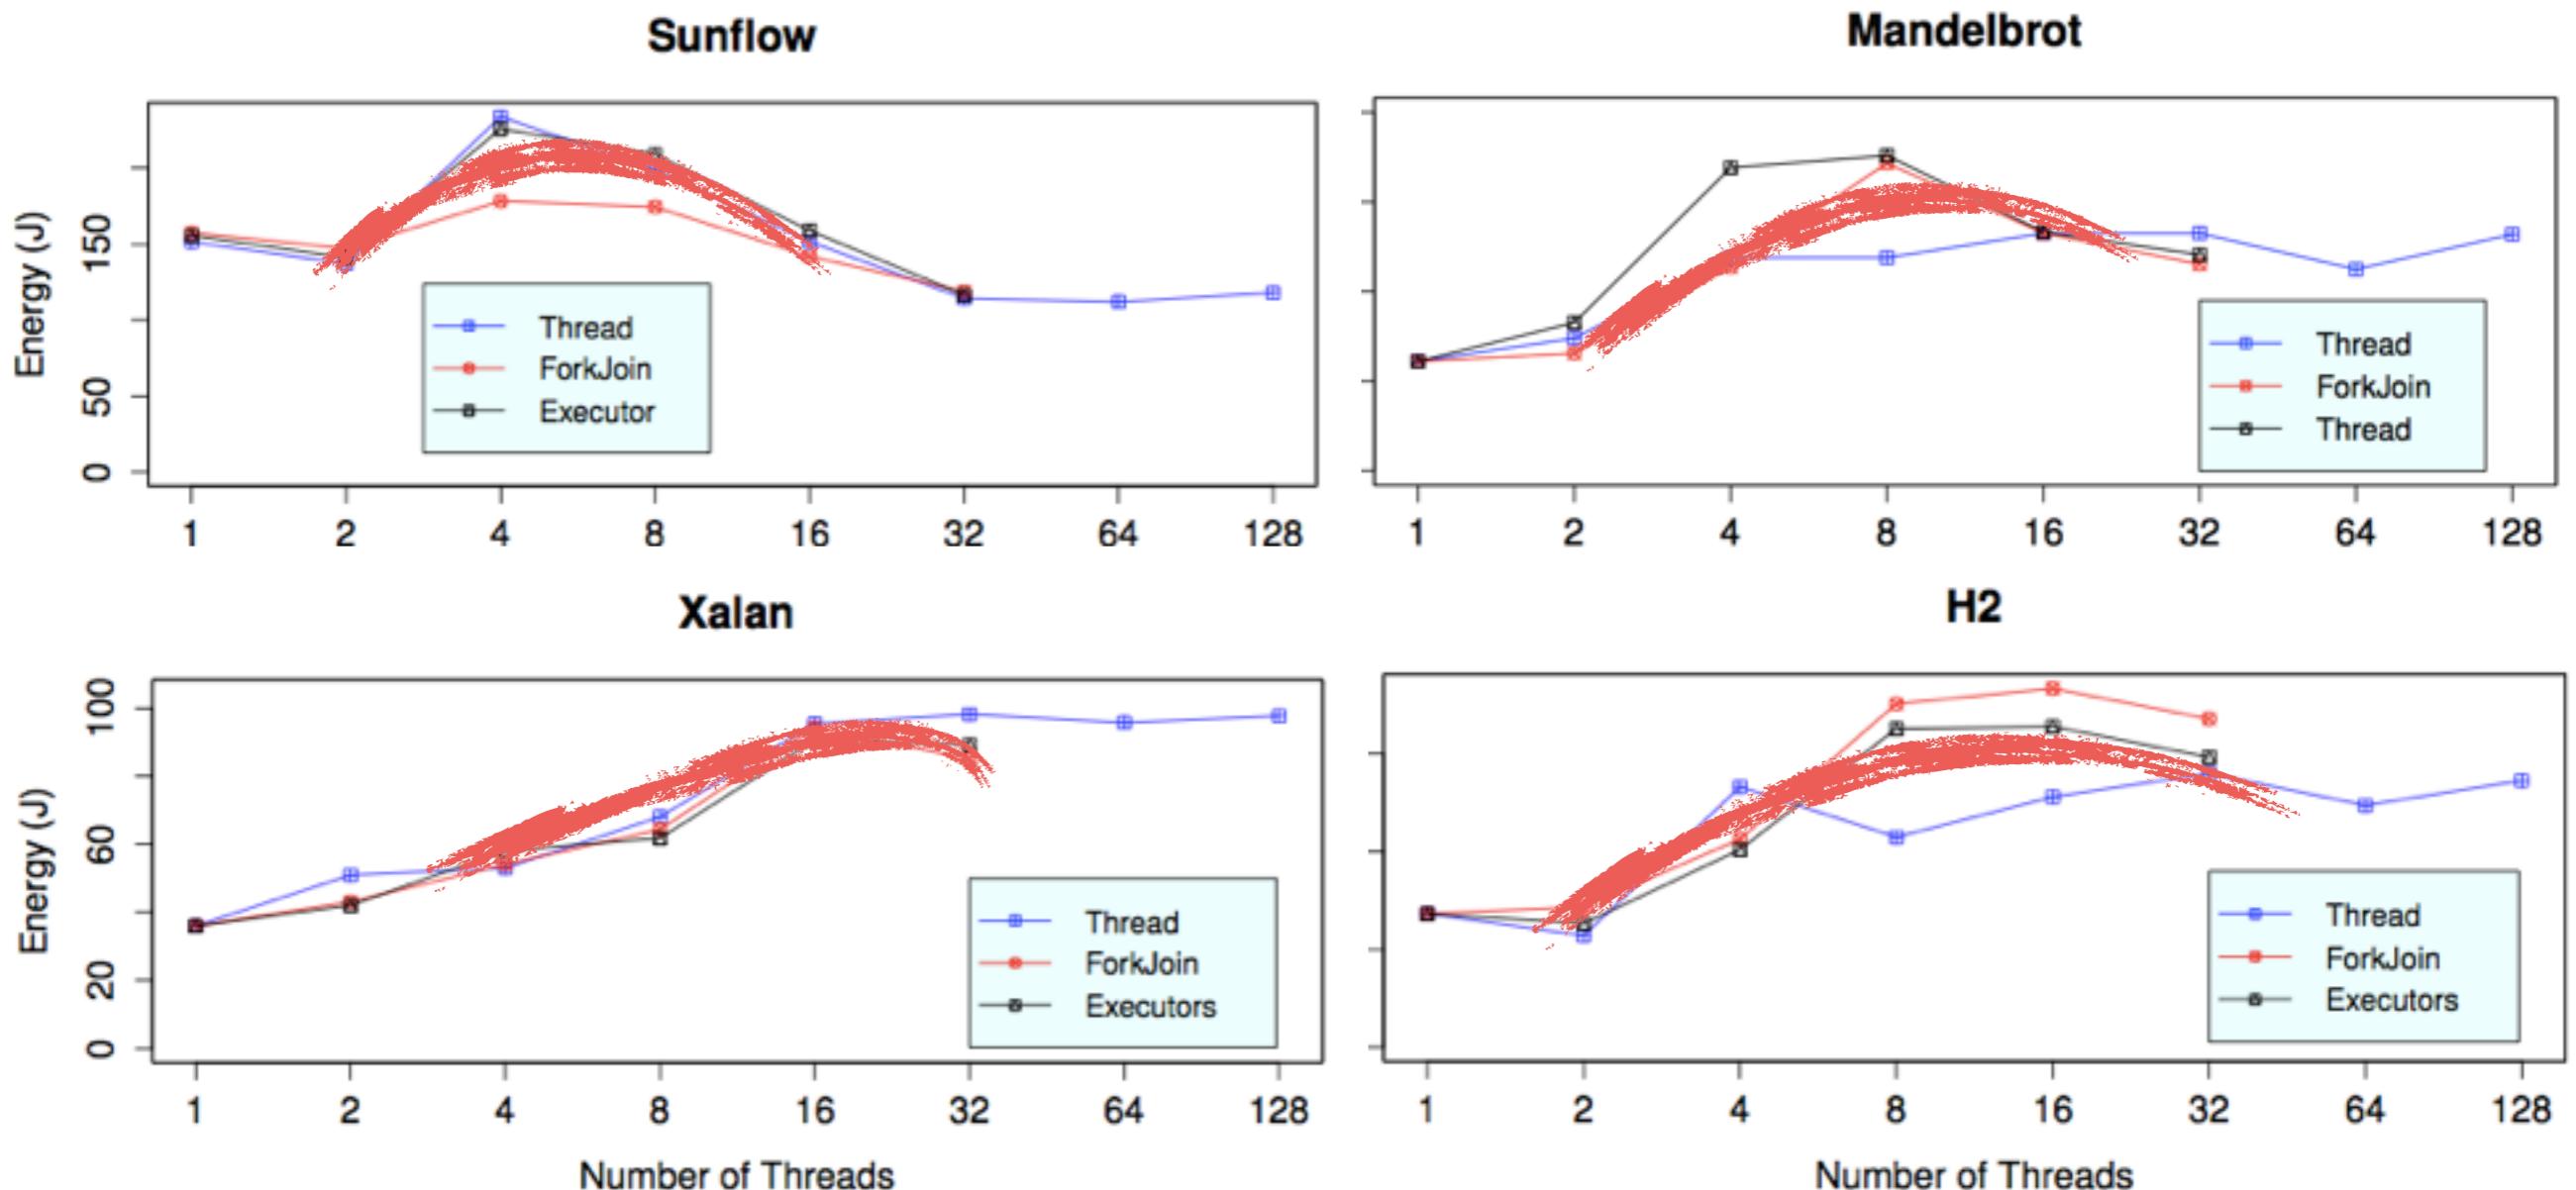
I Environment



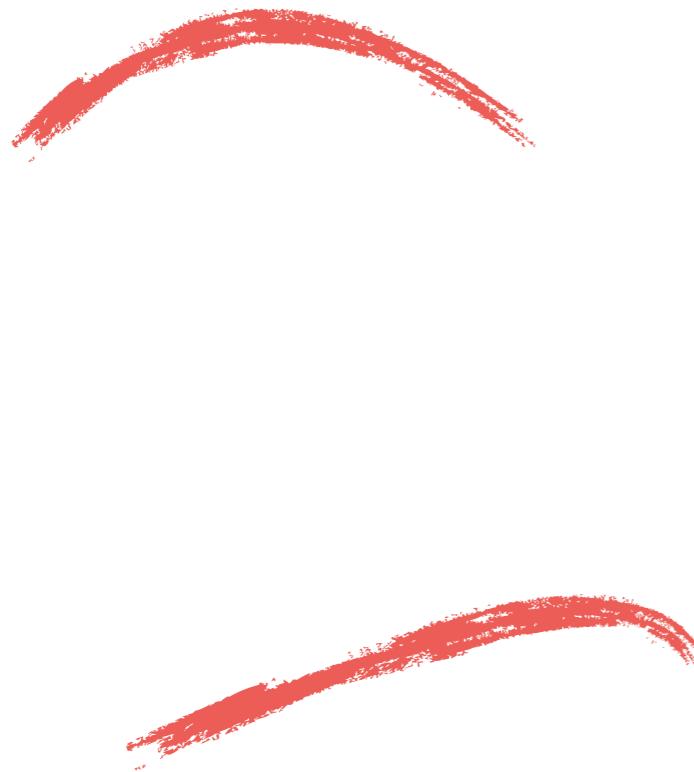
Energy Consumption When Varying the Number of Threads



The Λ Curve



The Λ Curve



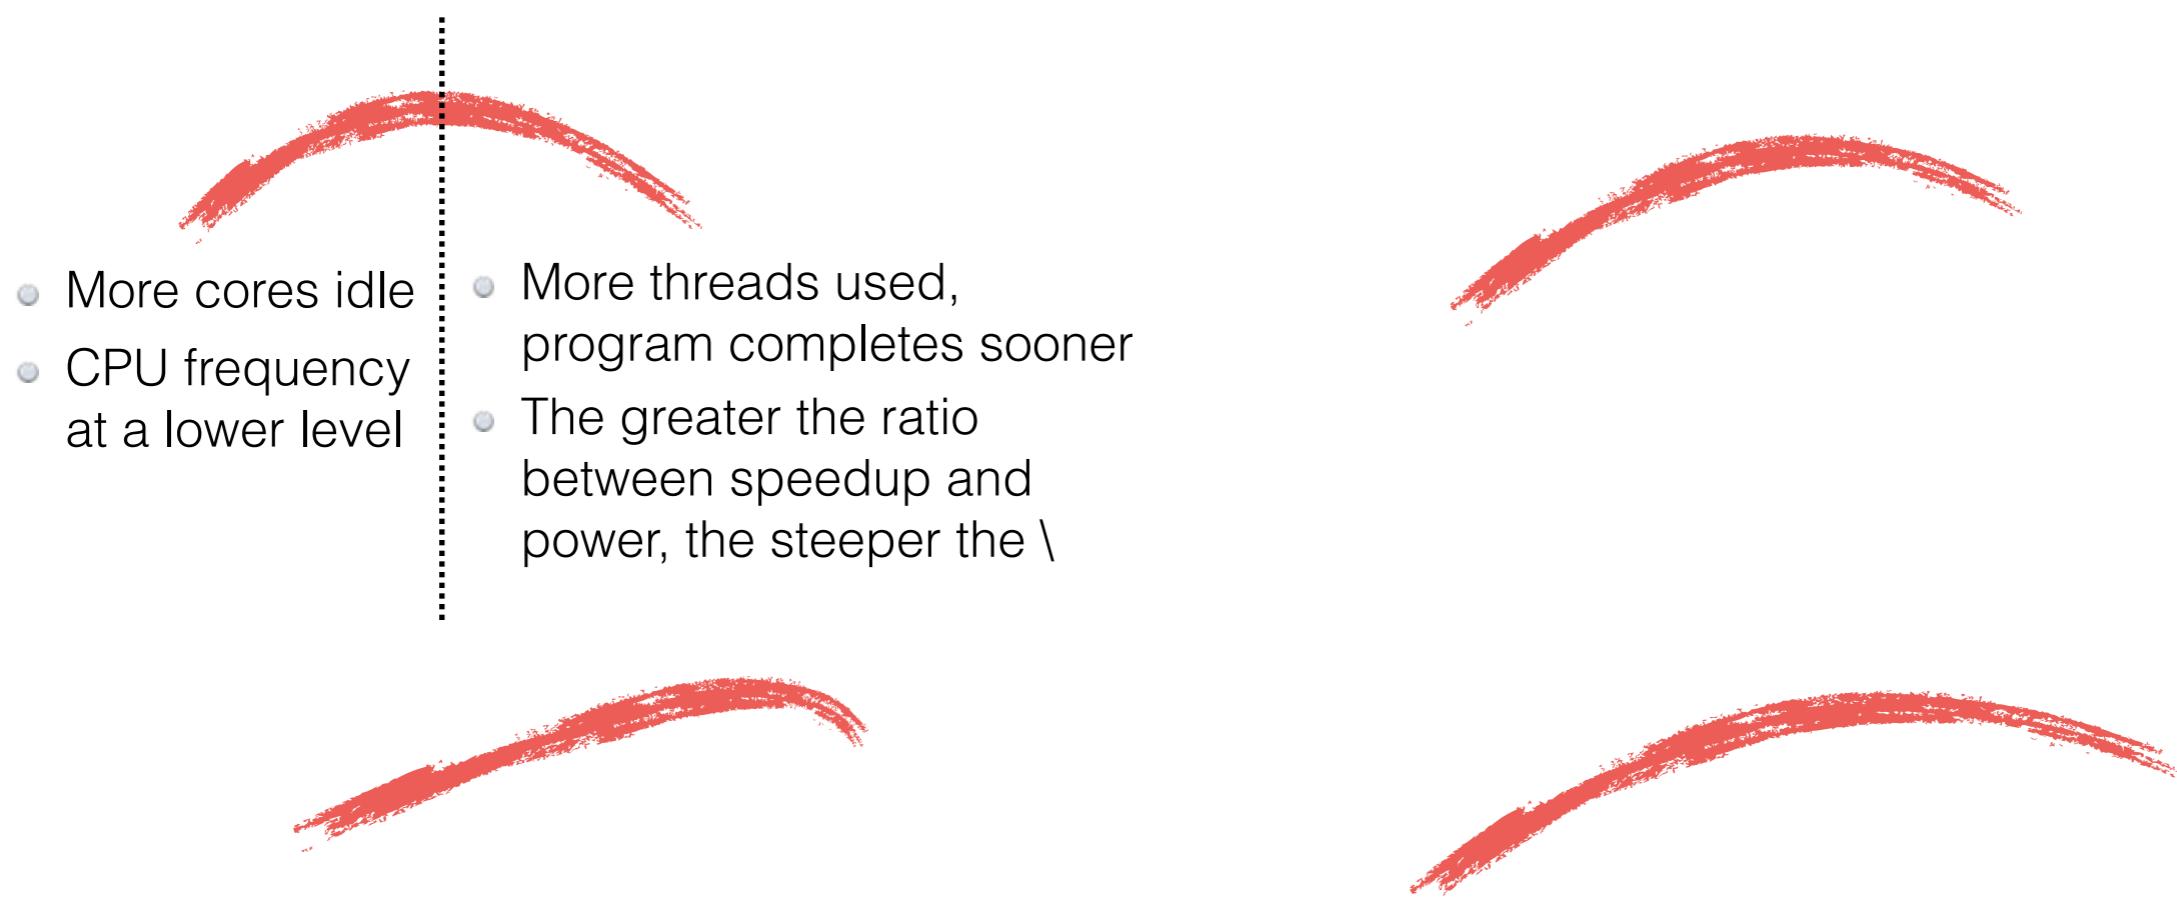
The Λ Curve



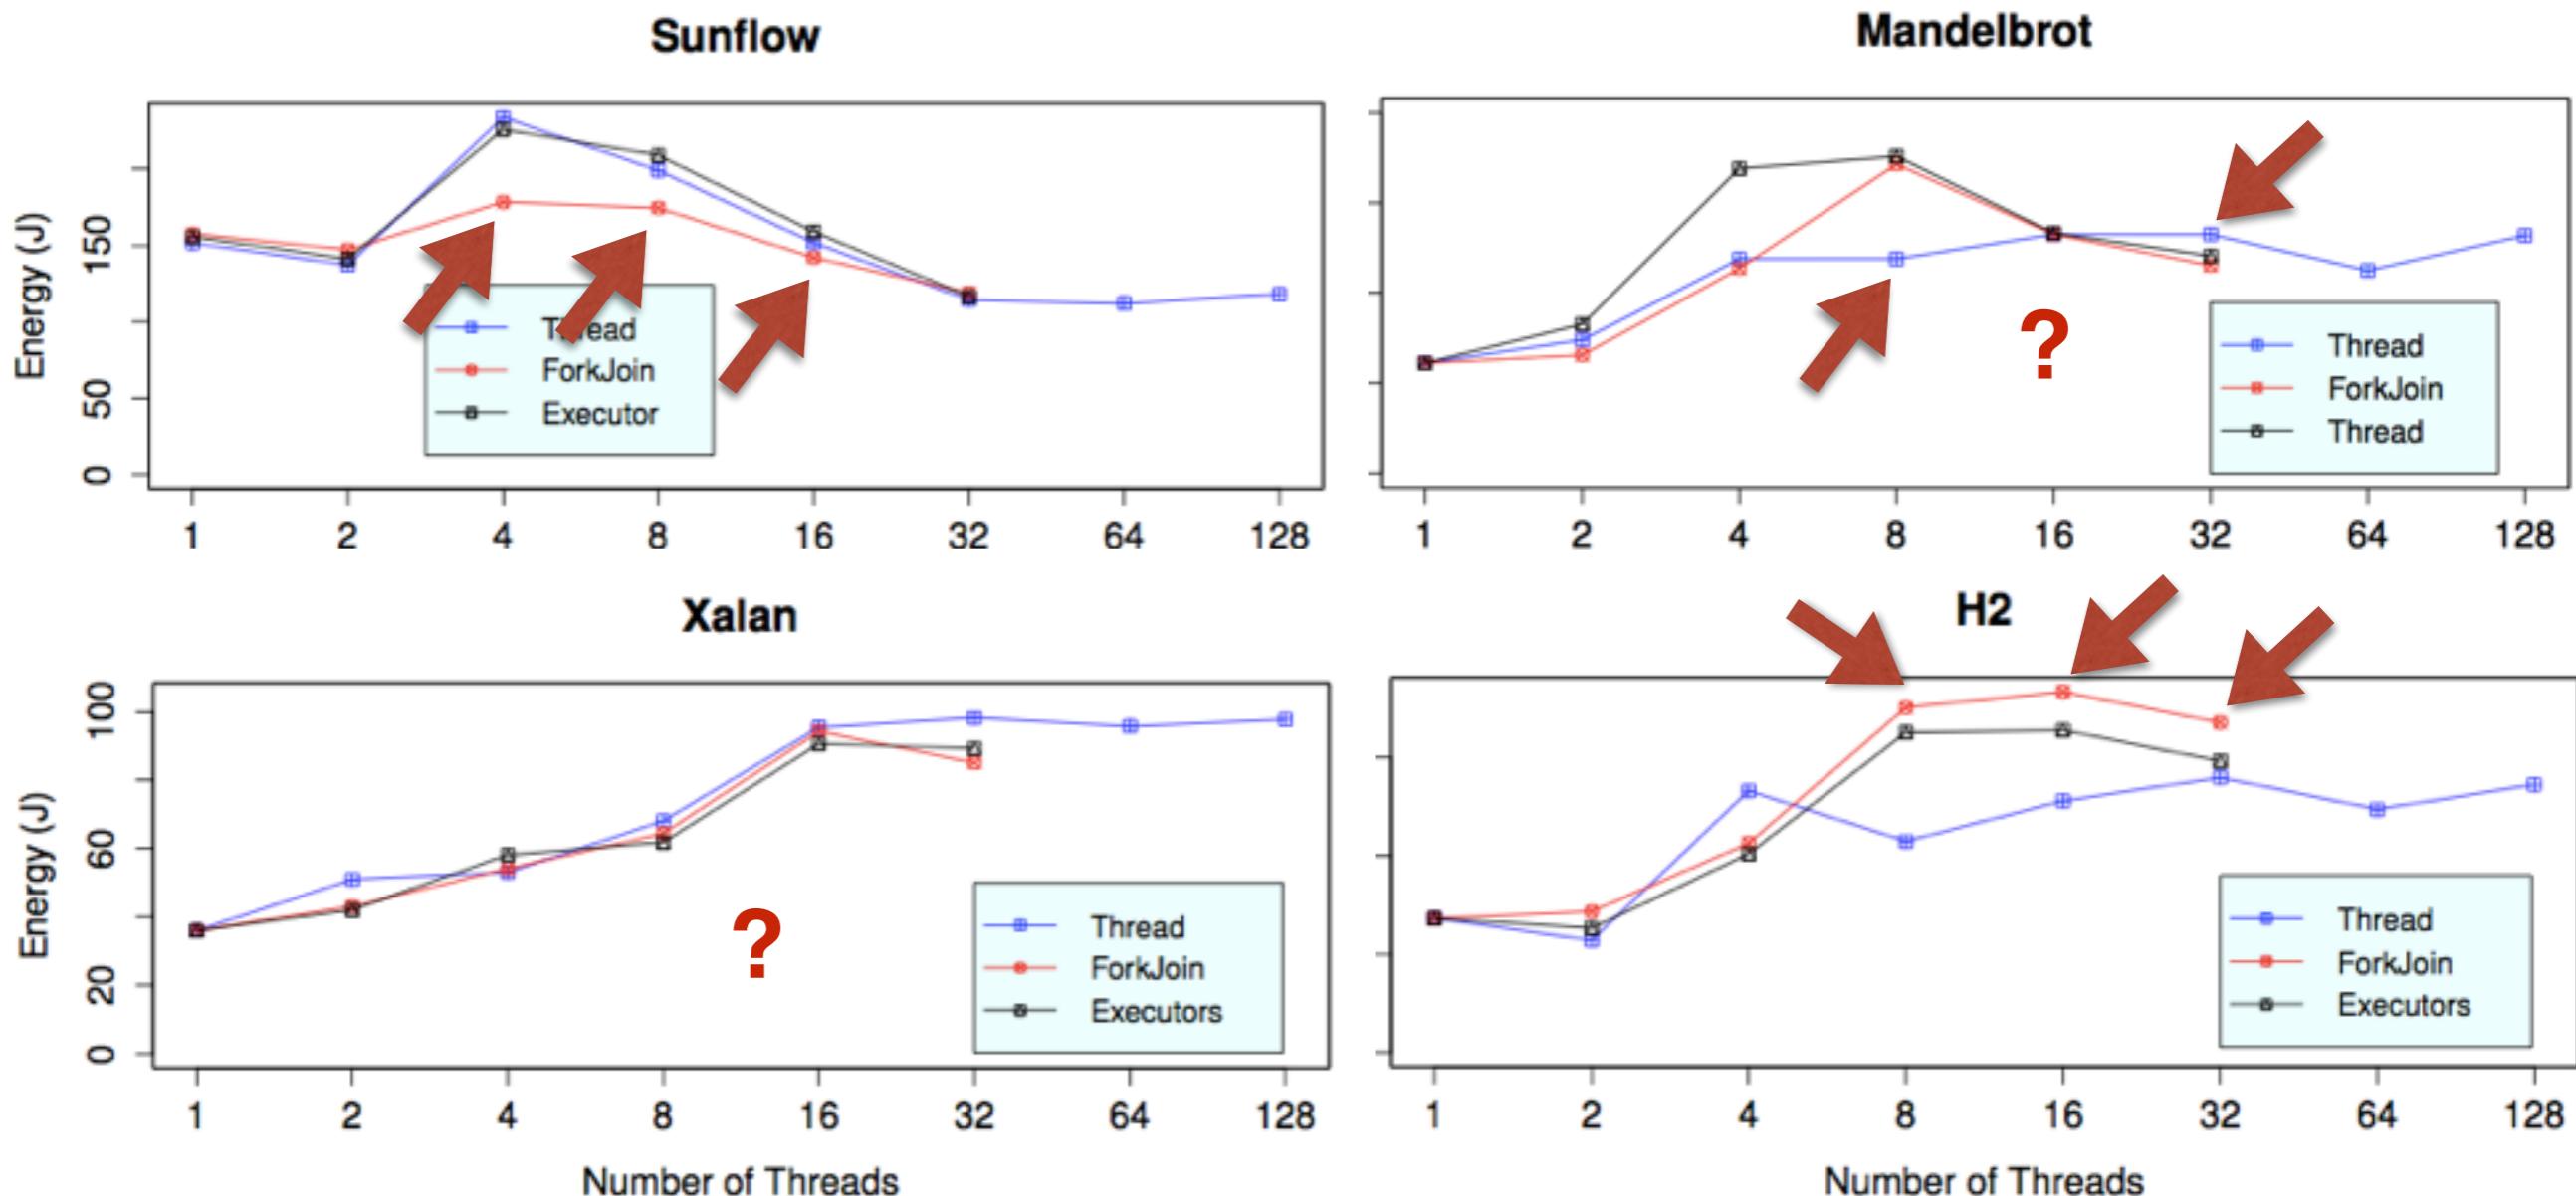
The Λ Curve



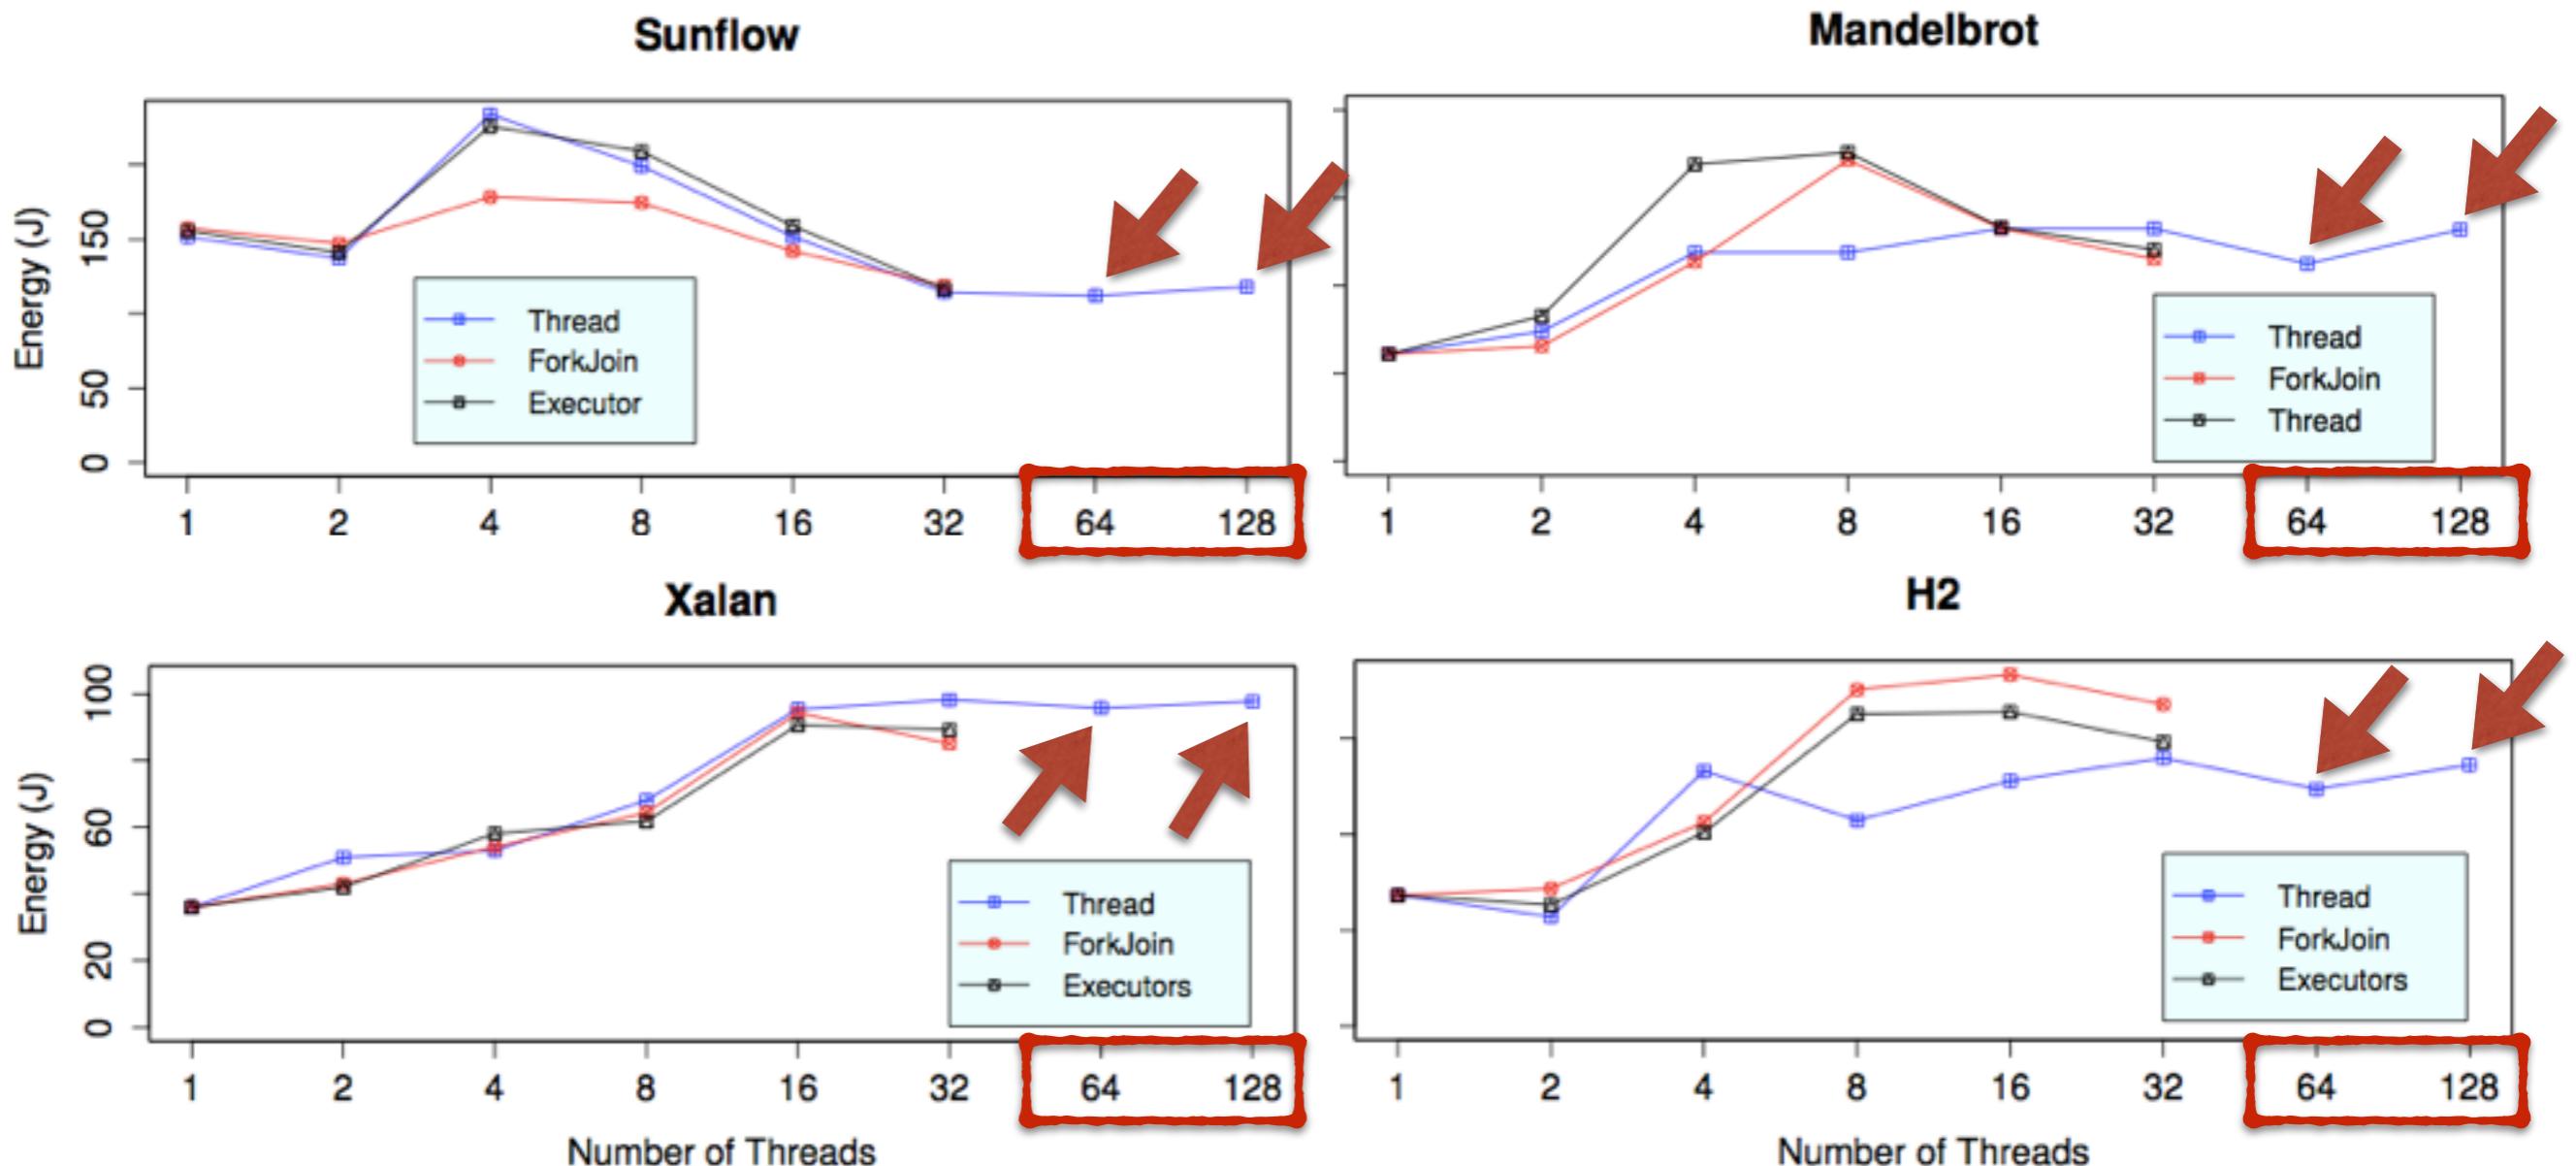
The Λ Curve

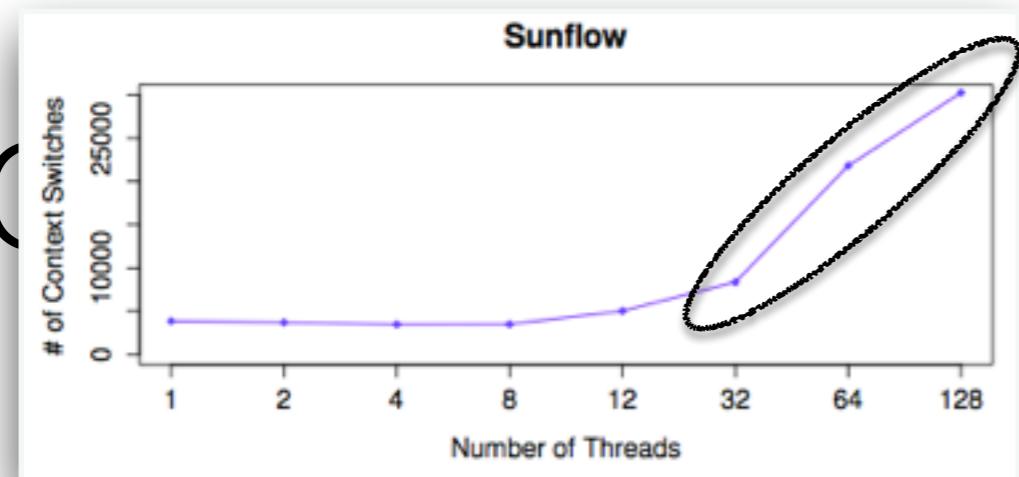


Which programming style should I use?



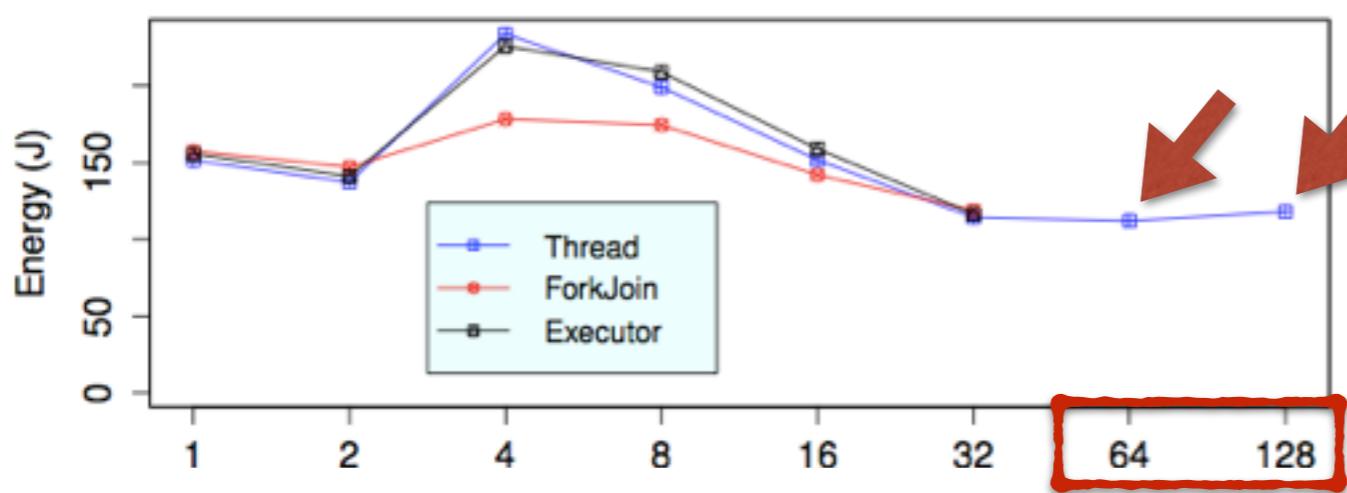
Overpopulating Cores with Threads



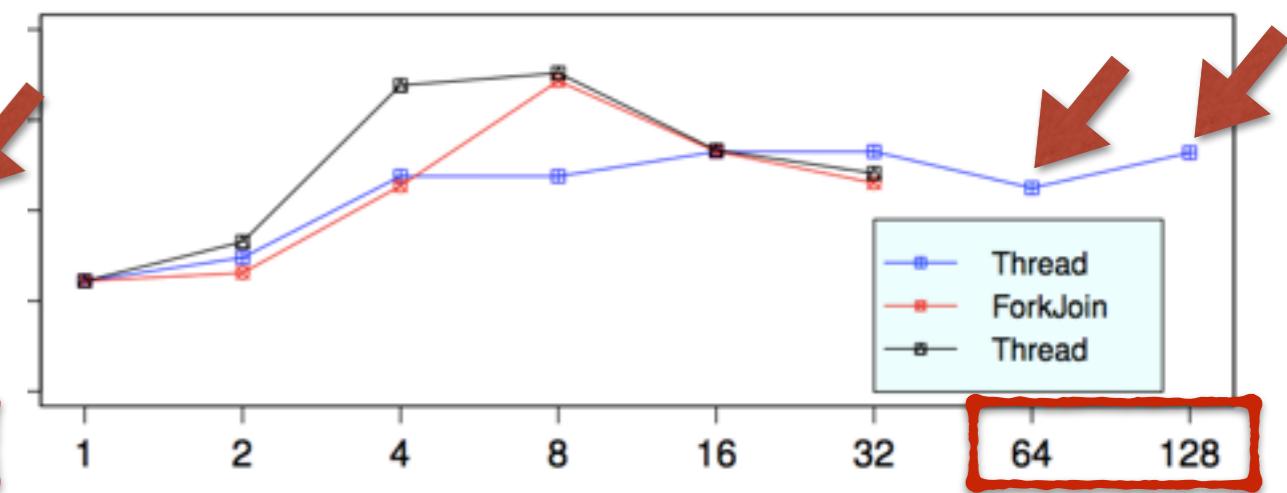


Cores with Threads

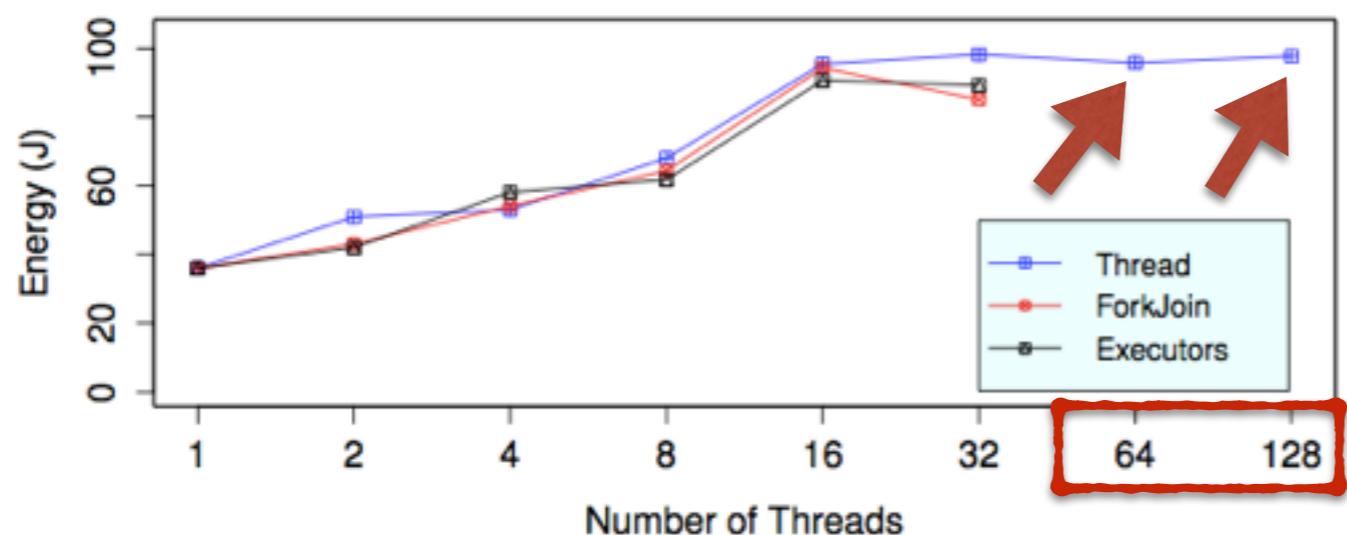
Sunflow



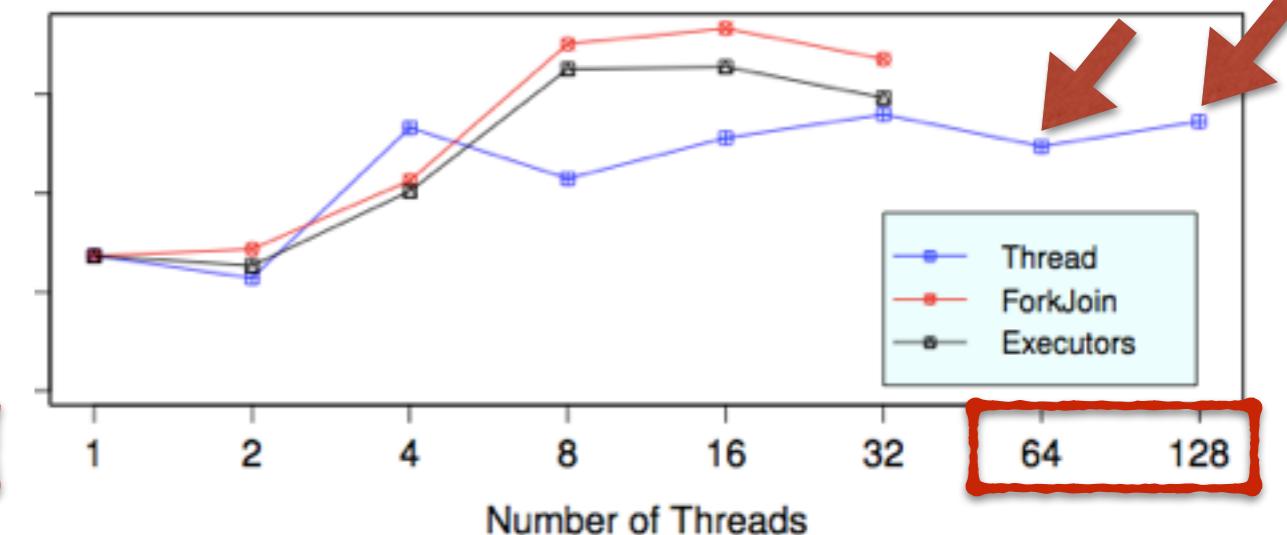
Mandelbrot



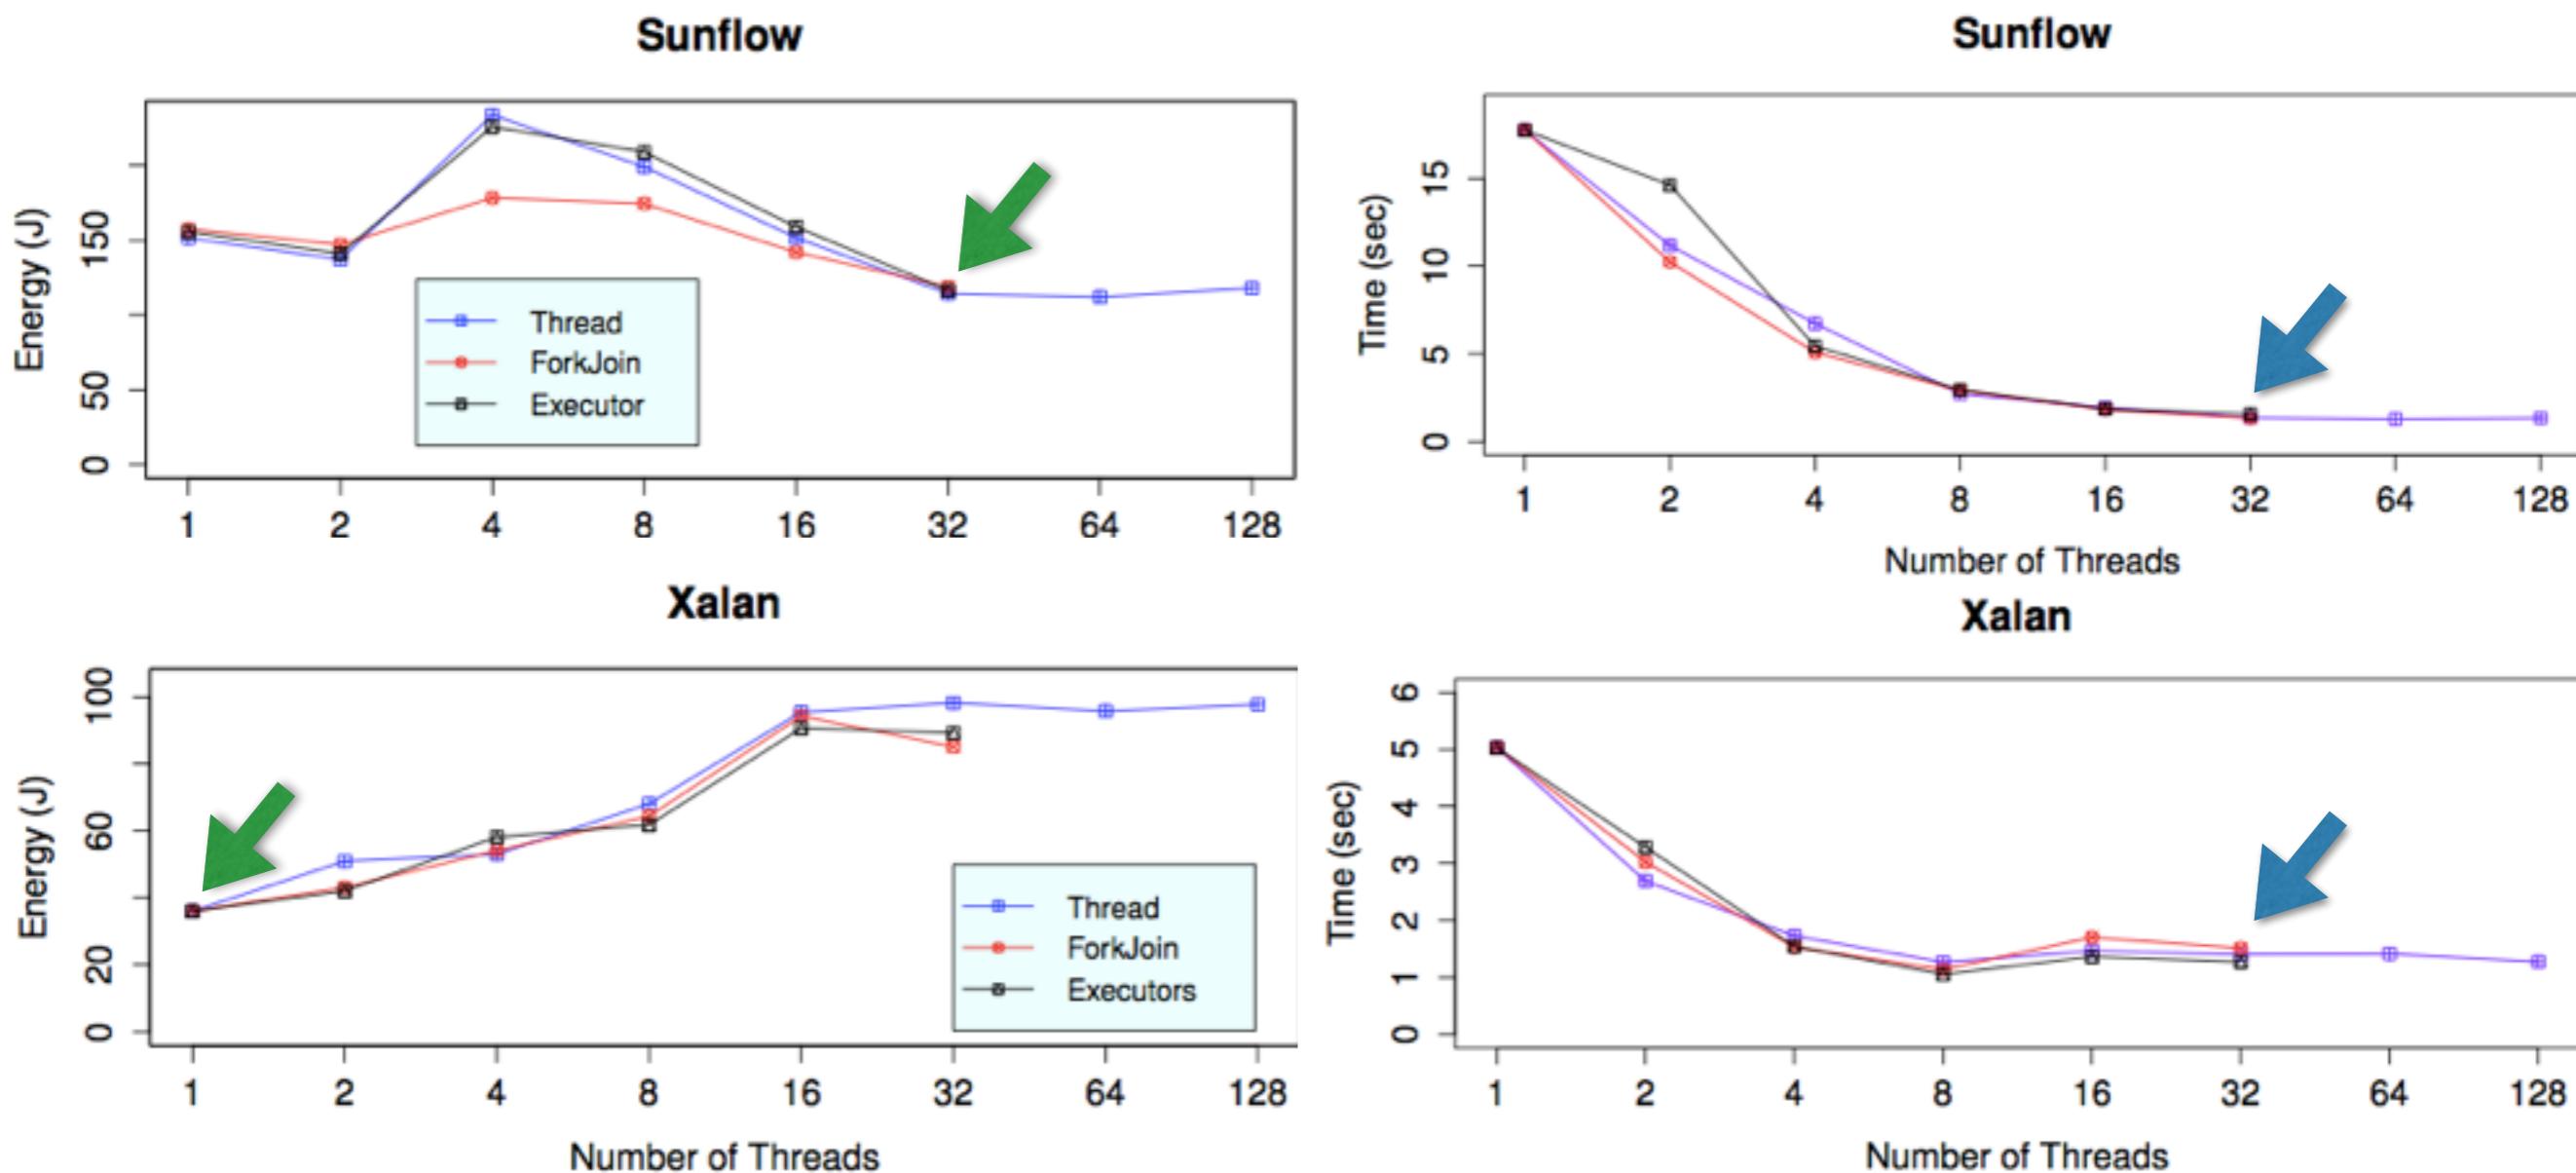
Xalan



H2



Faster ≠ Greener



Copying vs Sharing

```
import static Arrays.*;
class Task extends RecursiveAction{
    public Task (User[] u) {}
    protected void compute() {
        if (u.length < N) { local(u); }
        else {
            int split = u.length / 2;

            User[] u1 = copyOfRange(u, 0,
                split);
            User[] u2 = copyOfRange(u,
                split, u.length);

            invokeAll(new Task(u1),
                new Task(u2));
        }
    }
}
```

```
class Task extends RecursiveAction{
    public Task (User[] u, int from,
                int to) {}
    protected void compute() {
        if (to - from < N)
            local(u, from, to);
        else {
            int split = (from + to) / 2;

            invokeAll(
                new Task(u, from, split),
                new Task(u, from + split, to)
            );
        }
    }
}
```

Copying vs Sharing

```
import static Arrays.*;
class Task extends RecursiveAction{
    public Task (User[] u) {}
    protected void compute() {
        if (u.length < N) { local(u); }
        else {
            int split = u.length / 2;

            User[] u1 = copyOfRange(u, 0,
                split);
            User[] u2 = copyOfRange(u,
                split, u.length);

            invokeAll(new Task(u1),
                new Task(u2));
        }
    }
}
```

```
class Task extends RecursiveAction{
    public Task (User[] u, int from,
                int to) {}
    protected void compute() {
        if (to - from < N)
            local(u, from, to);
        else {
            int split = (from + to) / 2;

            invokeAll(
                new Task(u, from, split),
                new Task(u, from + split, to)
            );
        }
    }
}
```

Copying

Copying vs Sharing

```
import static Arrays.*;
class Task extends RecursiveAction{
    public Task (User[] u) {}
    protected void compute() {
        if (u.length < N) { local(u); }
        else {
            int split = u.length / 2;

            User[] u1 = copyOfRange(u, 0,
                split);
            User[] u2 = copyOfRange(u,
                split, u.length);

            invokeAll(new Task(u1),
                new Task(u2));
        }
    }
}
```

```
class Task extends RecursiveAction{
    public Task (User[] u, int from,
                int to) {}
    protected void compute() {
        if (to - from < N)
            local(u, from, to);
        else {
            int split = (from + to) / 2;

            invokeAll(
                new Task(u, from, split),
                new Task(u, from + split, to)
            );
        }
    }
}
```

Copying

Sharing

Copying vs Sharing

```
import static Arrays.*;
class Task extends RecursiveAction{
    public Task (User[] u) {}
    protected void compute() {
        if (u.length < N) { local(u); }
        else {
            int split = u.length / 2;

            User[] u1 = copyOfRange(u, 0,
                split);
            User[] u2 = copyOfRange(u,
                split, u.length);

            invokeAll(new Task(u1),
                new Task(u2));
        }
    }
}
```

```
class Task extends RecursiveAction{
    public Task (User[] u, int from,
                int to) {}
    protected void compute() {
        if (to - from < N)
            local(u, from, to);
        else {
            int split = (from + to) / 2;

            invokeAll(
                new Task(u, from, split),
                new Task(u, from + split, to)
            );
        }
    }
}
```

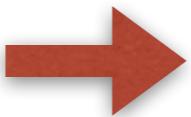
Copying

Sharing

±15% of energy savings!

Copy-Fork

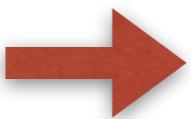
```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



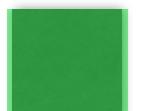
```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```

Copy-Fork

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```



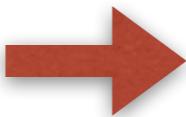
Before



After

Copy-Fork

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```

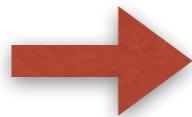
█ Before

█ After

Copy/Fork/Copy/Fork/...

Copy-Fork

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```

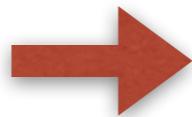
█ Before

█ After

Copy/Fork/Copy/Fork/...

Copy-Fork

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```



Before



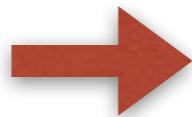
After

Copy./Copy/Fork.../Fork

Copy/Fork/Copy/Fork/...

Copy-Fork

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```



```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>(size);  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```



Before



After

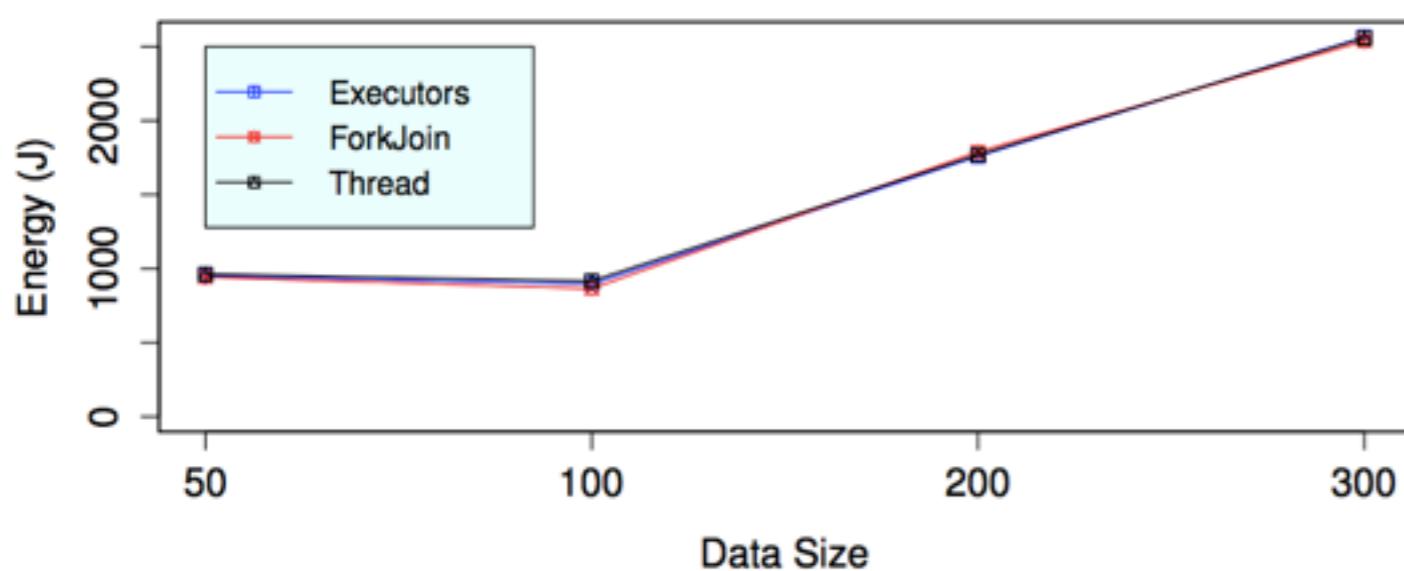
Copy./.Copy/Fork/.../Fork

±10% of energy
savings!

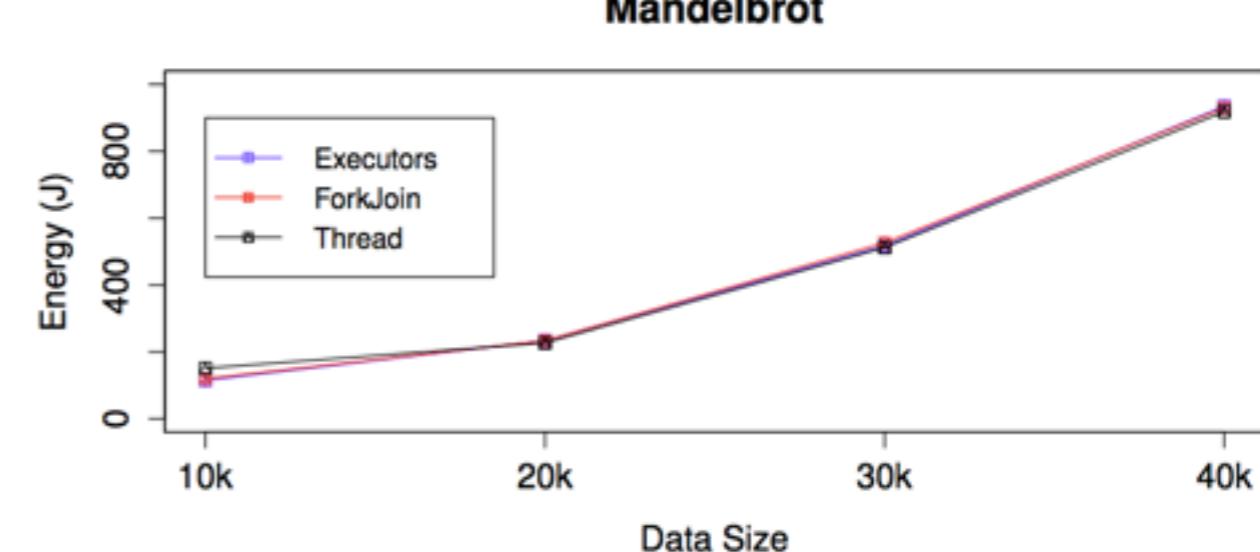
Copy/Fork/Copy/Fork/...

Data Size

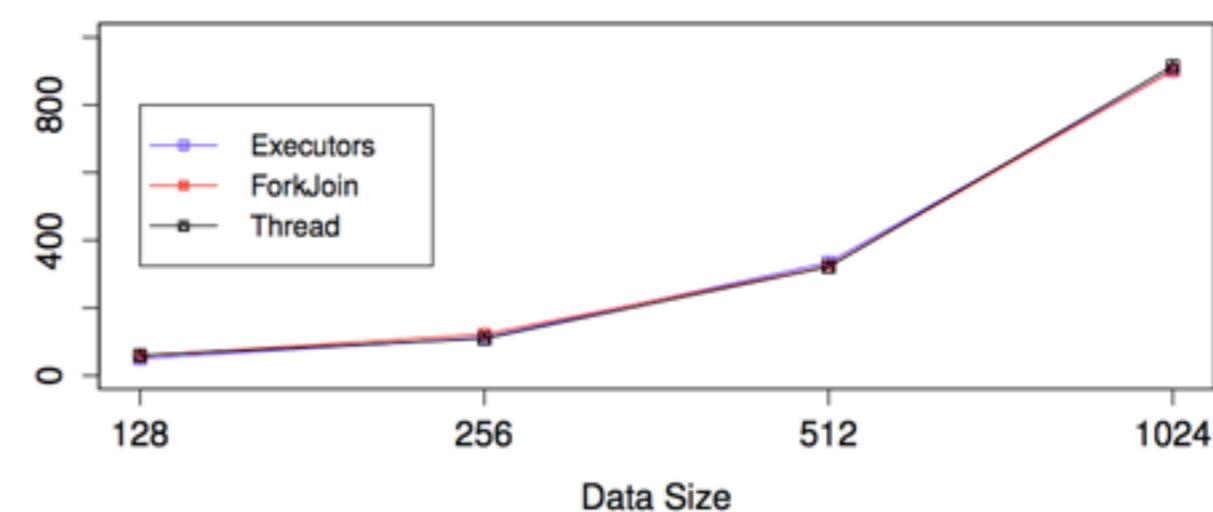
Xalan



Mandelbrot

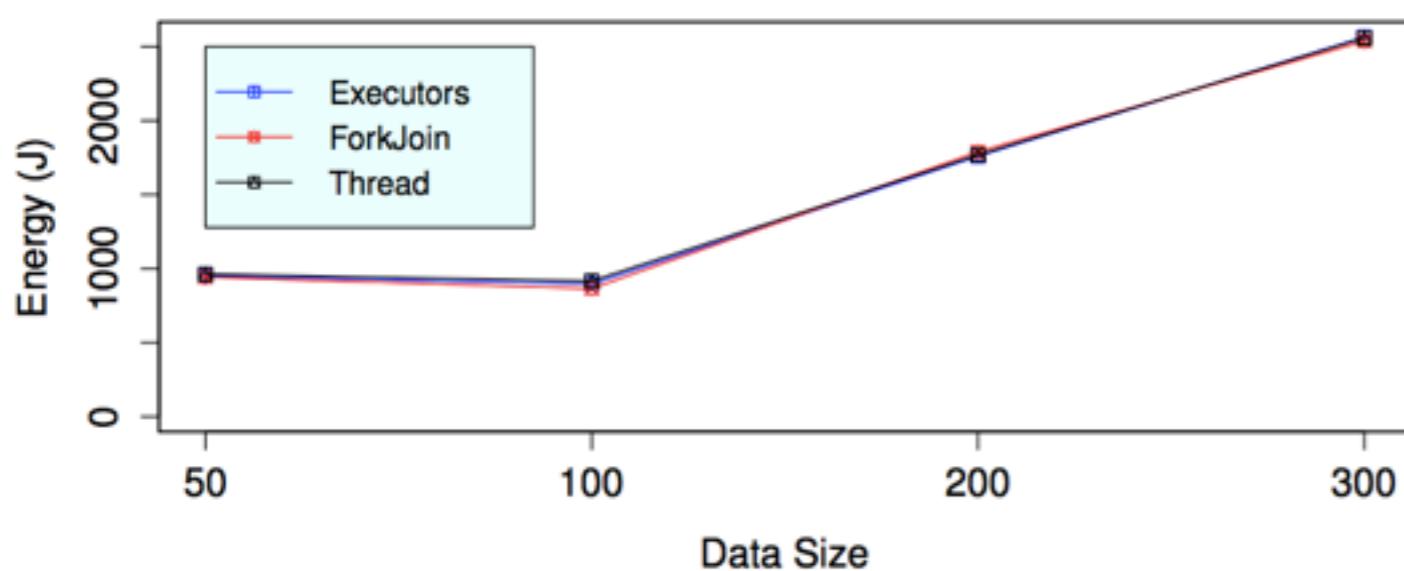


Sunflow

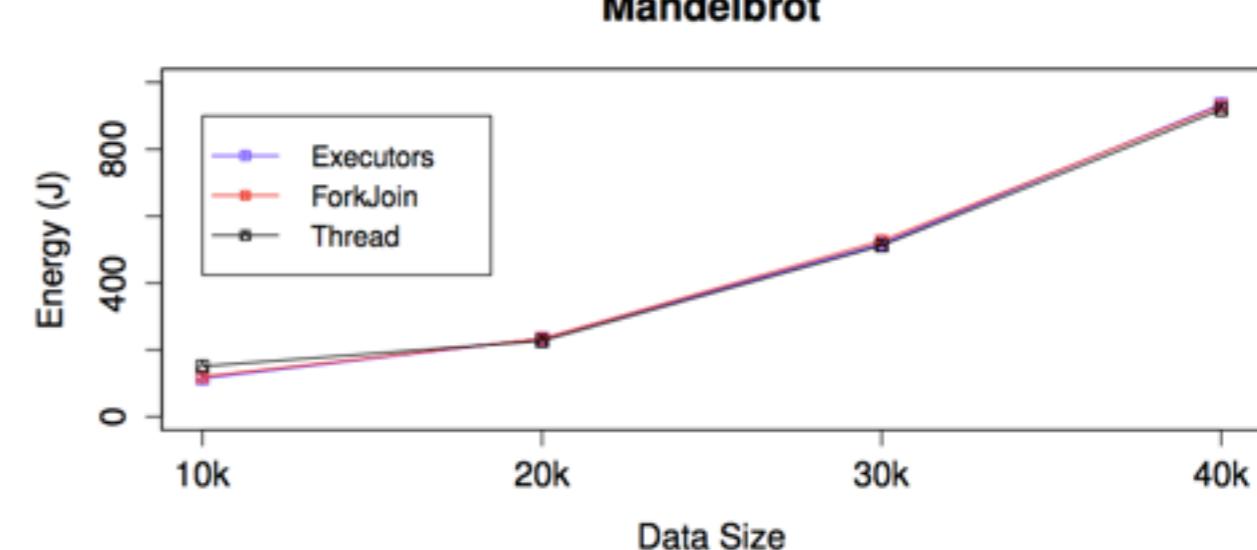


Data Size

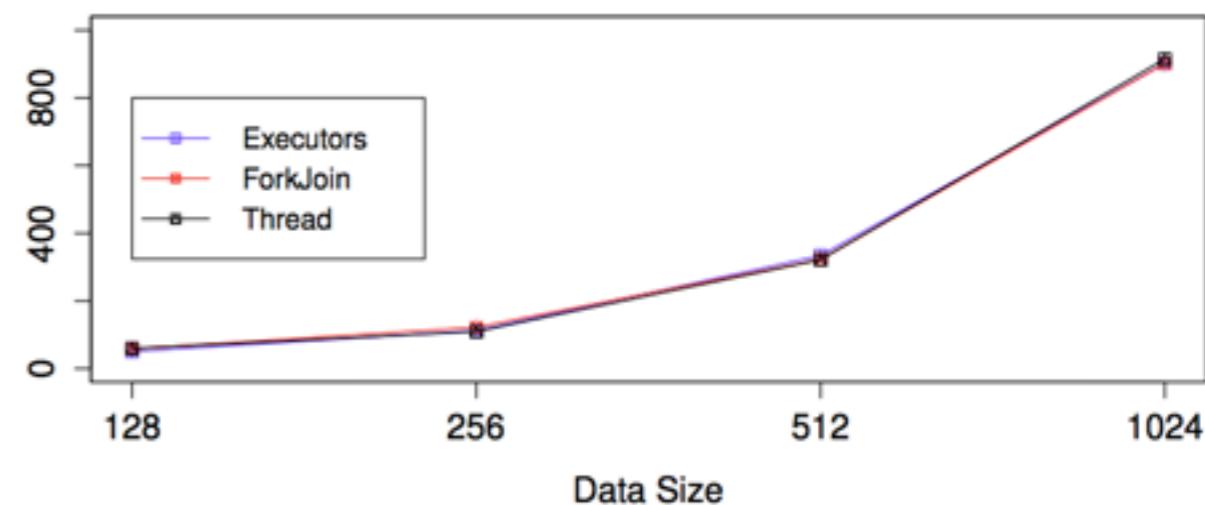
Xalan



Mandelbrot

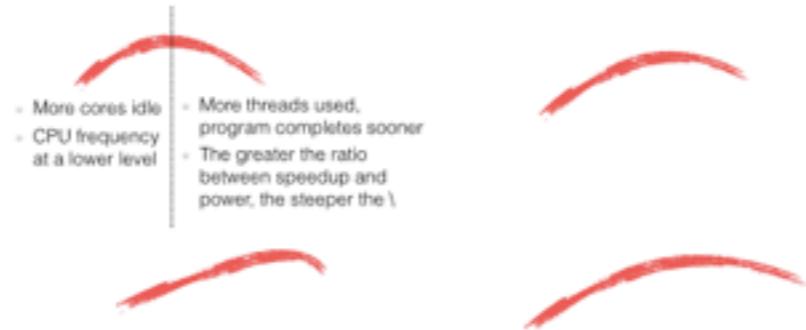


Sunflow



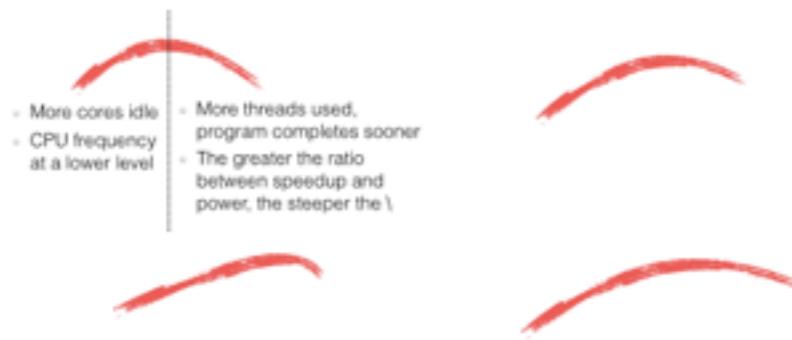
For most of the cases, energy consumption is linear to data size!

The Λ Curve



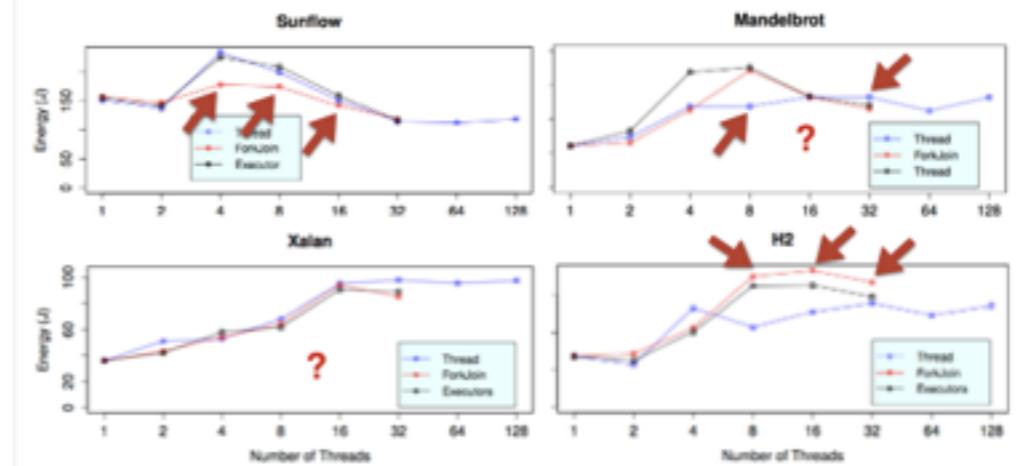
13

The Λ Curve



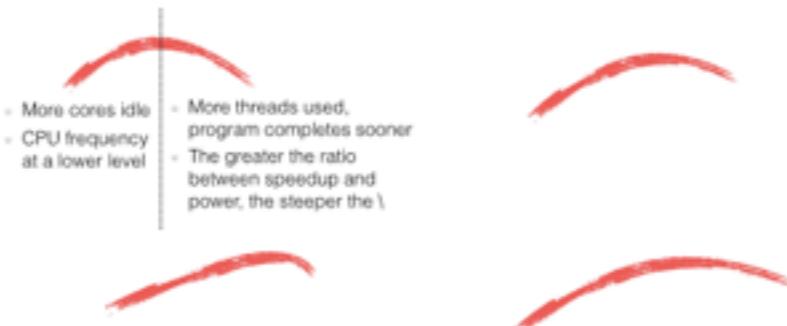
13

Which programming style should I use?



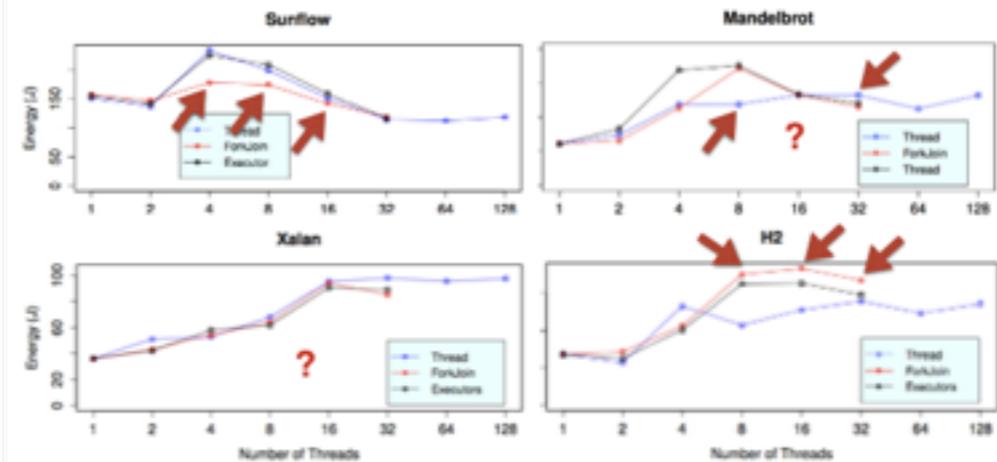
14

The Λ Curve



13

Which programming style should I use?



14

Copy-Fork Patterns

```
public void compute() {  
    ...  
    NQueensSolver[] tasks = new NQueensSolver[size];  
    for (int i = 0; i < tasks.length; i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks[i] = new NQueensSolver(newElements);  
        tasks[i].fork();  
    }  
    ...  
    for (int i = 0; i < tasks.length; i++) {  
        if(tasks[i] != null) tasks[i].join();  
    }  
}
```

→

```
public void compute() {  
    ...  
    List<NQueensSolver> tasks = new ArrayList<>{  
        size};  
    for (int i = 0; i < tasks.size(); i++) {  
        int[] newElements = new int[depth + 1];  
        System.arraycopy(currentElements, 0,  
                         newElements, 0, depth);  
        tasks.add(new NQueensSolver(newElements));  
    }  
    ...  
    invokeAll(tasks);  
}
```

Copy
Fork/Join

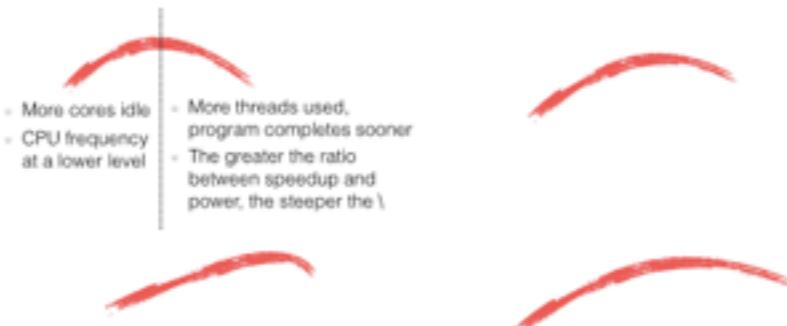
Copy/../Copy/Compute/Fork/.../Fork

Copy/Fork/Copy/Fork/.../Compute

±10% of energy
savings!

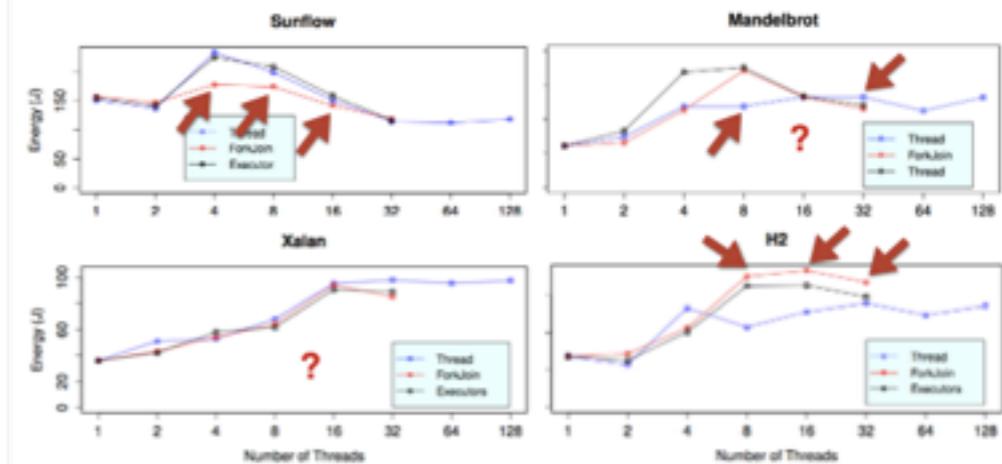
29

The Λ Curve



13

Which programming style should I use?



14

Copy-Fork Patterns

```
public void compute() {
    ...
    NQueensSolver[] tasks = new NQueensSolver[size];
    for (int i = 0; i < tasks.length; i++) {
        int[] newElements = new int[depth + 1];
        System.arraycopy(currentElements, 0,
                         newElements, 0, depth);
        tasks[i] = new NQueensSolver(newElements);
        tasks[i].fork();
    }
    ...
    for (int i = 0; i < tasks.length; i++) {
        if(tasks[i] != null) tasks[i].join();
    }
}
```

Copy

```
public void compute() {
    ...
    List<NQueensSolver> tasks = new ArrayList<>(
        size);
    for (int i = 0; i < tasks.size(); i++) {
        int[] newElements = new int[depth + 1];
        System.arraycopy(currentElements, 0,
                         newElements, 0, depth);
        tasks.add(new NQueensSolver(newElements));
    }
    ...
    invokeAll(tasks);
}
```

Fork/Join

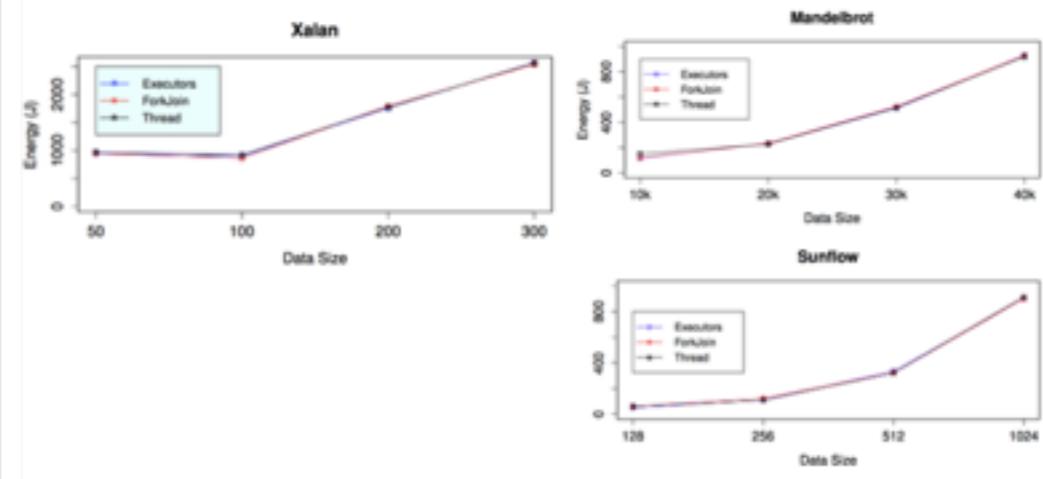
Copy/../Copy/Compute/Fork/.../Fork

Copy/Fork/Copy/Fork/.../Compute

±10% of energy savings!

29

Data Size



32

Understanding Energy Behaviors of Thread Management Constructs



Gustavo Pinto¹



Fernando Castor¹



David Liu²

{ghlp, castor}@cin.ufpe.br¹
davidL@binghamton.edu²



BINGHAMTON
UNIVERSITY
STATE UNIVERSITY OF NEW YORK

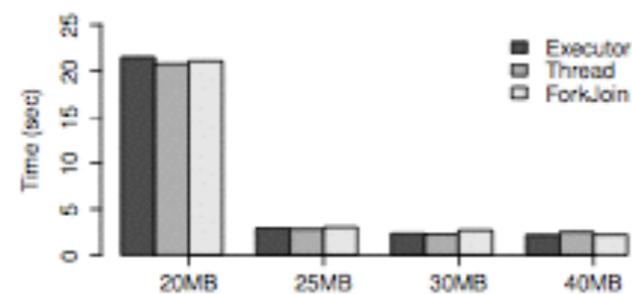
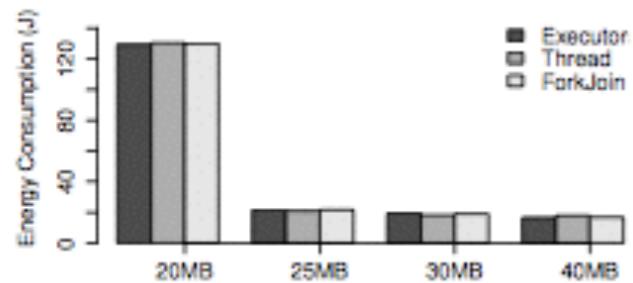


Figure 17. Heap Size Effect (sunflow, 32 threads, 256 tasks, 256 as image data size)

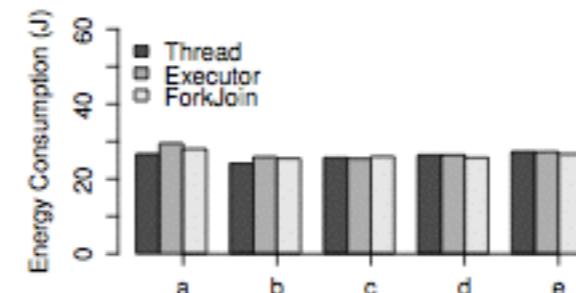
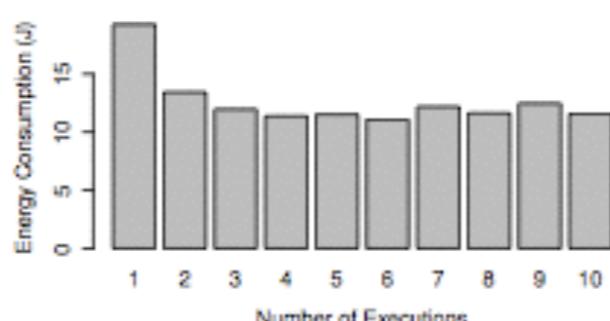
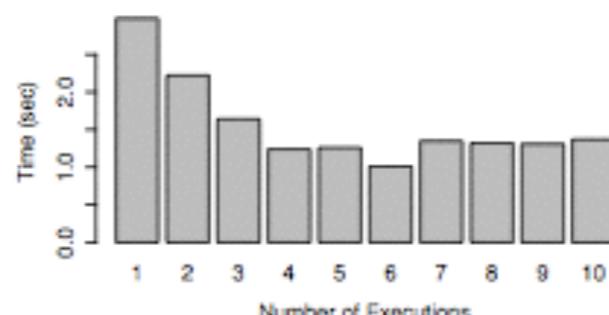


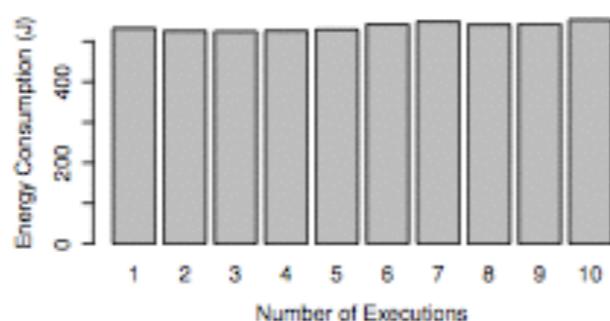
Figure 18. GC Effect (xalan, 32 threads, 64 tasks, 300 transformation files. GC strategies are: a: SerialGC, b: ParallelGC, c: ParallelOldGC, d: ConcMarkSweepGC, e: G1GC)



(a) Energy with JIT



(b) Time with JIT



(c) Energy without JIT



(d) Time without JIT

Figure 19. JIT Effect (sunflow, 32 threads, 256 tasks, 256 as image data size, 10 runs on X-axis)